

Tables and Priority Queues

Initially prepared by Dr. İlyas Çiçekli; improved by various Bilkent CS202 instructors.

Tables

- Appropriate for problems that must manage data by value.
- Some important operations of tables:
 - Inserting a data item containing the value x.
 - Delete a data item containing the value x.
 - Retrieve a data item containing the value x.
- Various table implementations are possible.
 - We have to analyze the possible implementations so that we can make an intelligent choice.
 - Some operations are implemented more efficiently in certain implementations.

An ordinary table of cities

<u>City</u>	<u>Country</u>	<u>Population</u>
Athens	Greece	2,500,000
Barcelona	Spain	1,800,000
Cairo	Egypt	9,500,000
London	England	9,400,000
New York	U.S.A.	7,300,000
Paris	France	2,200,000
Rome	Italy	2,800,000
Toronto	Canada	3,200,000
Venice	Italy	300,000

Table Operations

- Some of the table operations are possible:

- Create an empty table
- Destroy a table
- Determine whether a table is empty
- Determine the number of items in the table
- Insert a new item into a table
- Delete the item with a given search key
- Retrieve the item with a given search key
- Traverse the table

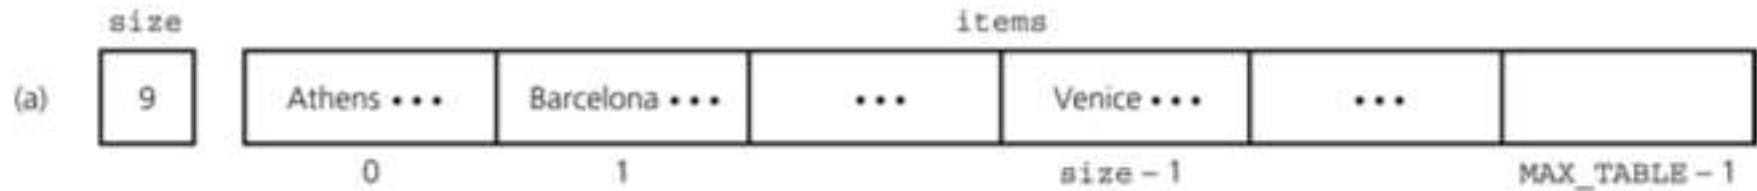
- The client may need a subset of these operations or require more
- Are keys in the table are unique?
 - We will assume that keys in our tables are unique.
 - But, some other tables allow duplicate keys.

Selecting an Implementation

- Since an array or a linked list represents items one after another, these implementations are called **linear**.
- There are four categories of linear implementations:
 - Unsorted, array based (an unsorted array)
 - Unsorted, pointer based (a simple linked list)
 - Sorted (by search key), array based (a sorted array)
 - Sorted (by search key), pointer based (a sorted linked list).
- We have also **nonlinear** implementations such as binary search trees.
 - Binary search tree implementation offers several advantages over linear implementations.

Sorted Linear Implementations

Array-based implementation

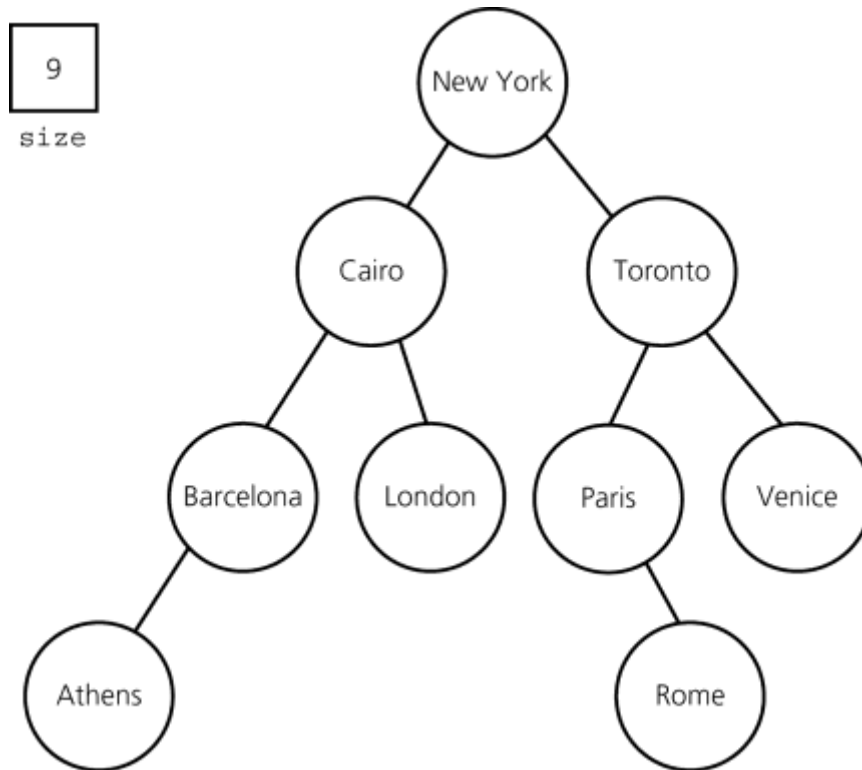


Pointer-based implementation



A Nonlinear Implementation

Binary search tree implementation



Which Implementation?

- It depends on our application.
- Answer the following questions before selecting an implementation.

1. What operations are needed?

- Our application may not need all operations.
- Some operations can be implemented more efficiently in one implementation, and some others in another implementation.

2. How often is each operation required?

- Some applications may require many occurrences of an operation, but other applications may not.
 - For example, some applications may perform **many retrievals**, but not so many insertions and deletions. On the other hand, other applications may perform **many insertions and deletions**.

How to Select an Implementation – Scenario A

- **Scenario A:** Let us assume that we have an application:
 - Inserts data items into a table.
 - After all data items are inserted, traverses this table in **no particular order**.
 - Does not perform any retrieval and deletion operations.
 - Which implementation is appropriate for this application?
 - Keeping the items in a sorted order provides no advantage for this application.
 - In fact, it will be more costly for this application.
- *Unsorted implementation is more appropriate.*

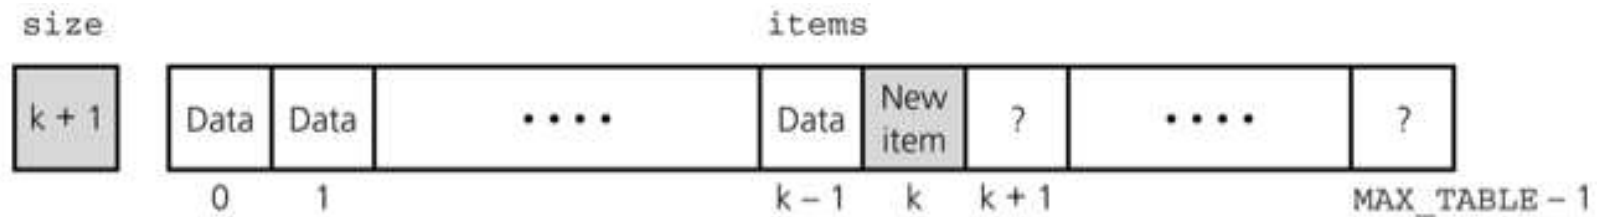
How to Select an Implementation – Scenario A

- Which unsorted implementation (array-based, pointer-based)?
 - Do we know the maximum size of the table?

- If we know the expected size is close to the maximum size of the table

→ *an array-based implementation is more appropriate*

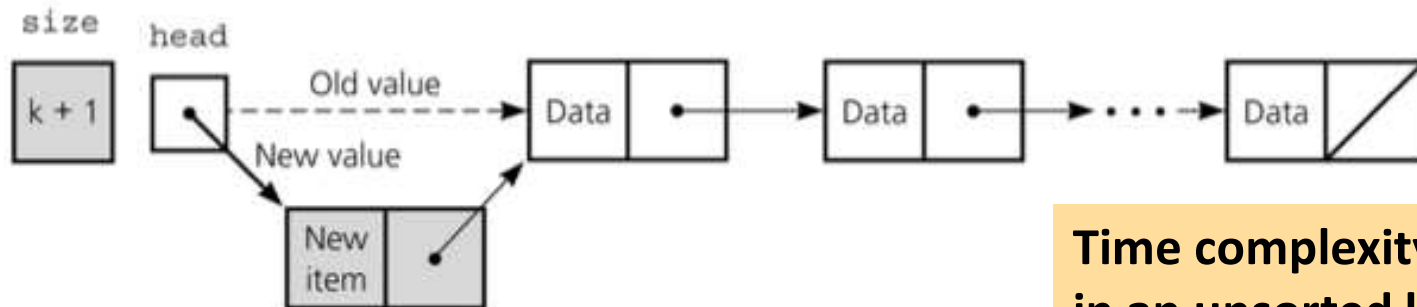
(because a pointer-based implementation uses extra space for pointers)



- Otherwise,

→ *a pointer-based implementation is more appropriate*

(because too many entries will be empty in an array-based implementation)



**Time complexity of insertion
in an unsorted list: $O(1)$**

How to Select an Implementation – Scenario B

- **Scenario B:** Let us assume that we have an application:
 - Performs many retrievals, but few insertions and deletions
 - E.g., a thesaurus (to look up synonyms of a word)
- For this application, *a sorted implementation is more appropriate*
 - We can use binary search to access data, if we have sorted data.
 - A sorted linked-list implementation is not appropriate since binary search is not practical with linked lists.
 - If we know the maximum size of the table
 - ➔ *a sorted array-based implementation is more appropriate for frequent retrievals.*
 - Otherwise
 - ➔ *a binary search tree implementation is more appropriate for frequent retrievals.*
(in fact, balanced binary search trees will be used)

How to Select an Implementation – Scenario C

- **Scenario C:** Let us assume that we have an application:
 - Performs many retrievals as well as many insertions and deletions.

? Sorted Array Implementation

- Retrievals are efficient.
- But insertions and deletions are not efficient.

shifting data

→ *a sorted array-based implementation is not appropriate for this application.*

? Sorted Linked List Implementation

- Retrievals, insertions, and deletions are not efficient.

→ *a sorted linked-list implementation is not appropriate for this application.*

? Binary Search Tree Implementation

- Retrieval, insertion, and deletion are efficient in the average case.

→ *a binary search tree implementation is appropriate for this application.*
(provided that the height of the BST is $O(\log n)$)

Which Implementation?

- Linear implementations of a table can be appropriate despite its difficulties.
 - Linear implementations are easy to understand, easy to implement.
 - For small tables, linear implementations can be appropriate.
 - For large tables, linear implementations may still be appropriate (e.g., for the case that has only insertions to an unsorted table--Scenario A)
- In general, a **binary search tree** implementation is a better choice.
 - Worst case: $O(n)$ for most table operations
 - Average case: $O(\log_2 n)$ for most table operations
- **Balanced binary search trees** increase the efficiency.

Which Implementation?

The average-case time complexities of the table operations

	<u>Insertion</u>	<u>Deletion</u>	<u>Retrieval</u>	<u>Traversal</u>
Unsorted array based	$O(1)$	$O(n)$	$O(n)$	$O(n)$
Unsorted pointer based	$O(1)$	$O(n)$	$O(n)$	$O(n)$
Sorted array based	$O(n)$	$O(n)$	$O(\log n)$	$O(n)$
Sorted pointer based	$O(n)$	$O(n)$	$O(n)$	$O(n)$
Binary search tree	$O(\log n)$	$O(\log n)$	$O(\log n)$	$O(n)$

Binary Search Tree Implementation – TableB.h

```
#include "BST.h"// Binary search tree operations
typedef TreeItemType TableItemType;

class Table {
public:
    Table();        // default constructor
    // copy constructor and destructor are supplied by the compiler

    bool tableIsEmpty() const;
    int  tableLength() const;
    void tableInsert(const TableItemType& newItem) throw(TableException);
    void tableDelete(KeyType searchKey) throw(TableException);
    void tableRetrieve(KeyType searchKey, TableItemType& tableItem) const
                                                throw(TableException);

    void traverseTable(FunctionType visit);

protected:
    void setSize(int newSize);
private:
    BinarySearchTree bst;                // BST that contains the table's items
    int size;                           // Number of items in the table
}
```

Binary Search Tree Implementation – tableInsert

```
#include "TableB.h"// header file

void Table::tableInsert(const TableItemType& newItem) throw(TableException) {
    try {
        bst.searchTreeInsert(newItem);
        ++size;
    }
    catch (TreeException e){
        throw TableException("Cannot insert item");
    }
}
```

The Priority Queue

Priority queue is a variation of the table.

- Each data item in a priority queue has a **priority value**.
- Using a priority queue we prioritize a list of tasks:
 - Job scheduling

Major operations:

- **Insert** an item with a priority value into its proper position in the priority queue.
- **Deletion is not the same** as the deletion in the table. We delete the item with the **highest priority**.

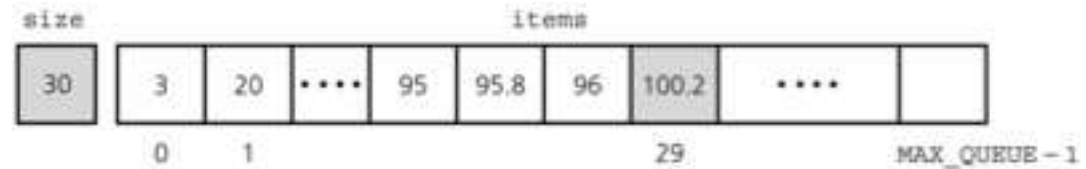
Priority Queue Operations

- create** – creates an empty priority queue.
- destroy** – destroys a priority queue.
- isEmpty** – determines whether a priority queue is empty or not.
- insert** – inserts a new item (with a priority value) into a priority queue.
- delete** – retrieves the item in a priority queue with the highest priority value, and deletes that item from the priority queue.

Which Implementations?

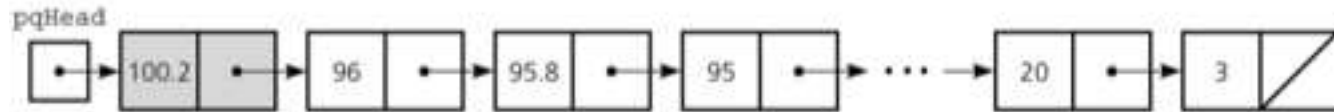
1. Array-based implementation

- Insertion will be $O(n)$



2. Linked-list implementation

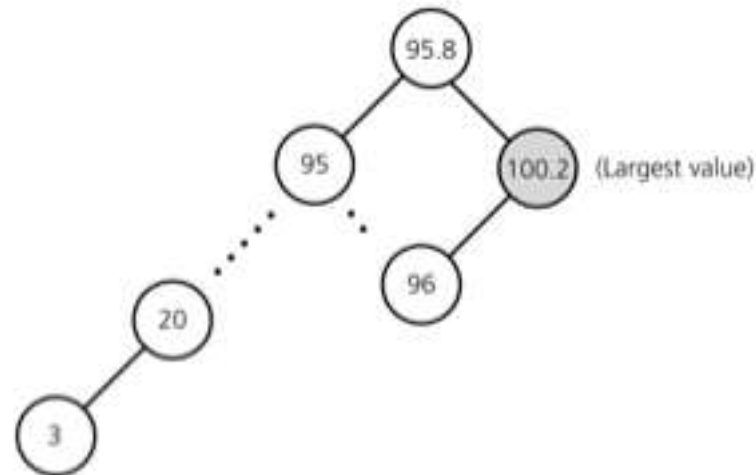
- Insertion will be $O(n)$



3. BST implementation

- Insertion is $O(\log_2 n)$ in average but $O(n)$ in the worst case.

We need a balanced BST so that we can get better performance [$O(\log n)$ in the worst case] → HEAP

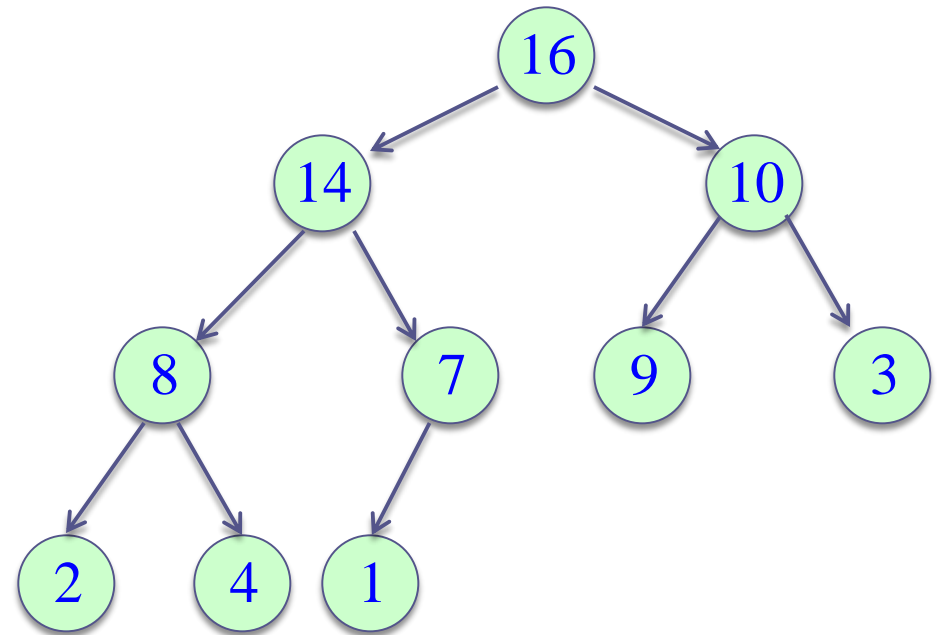
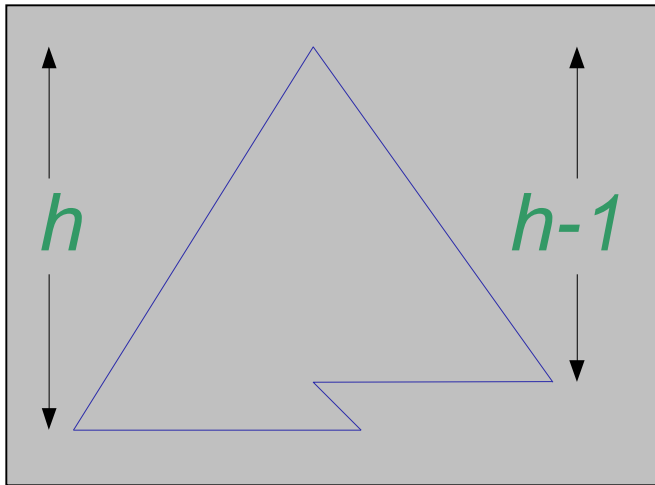


Heaps

Definition: A **heap** is a complete binary tree such that

- It is empty, or
 - Its root contains a search key greater than or equal to the search key in each of its children, and each of its children is also a heap.
-
- Since the root contains the item with the largest search key, heap in this definition is also known as **maxheap**.
 - On the other hand, a heap which places the smallest search key in its root is known as **minheap**.
 - We will talk about maxheap as heap in the rest of our discussions.

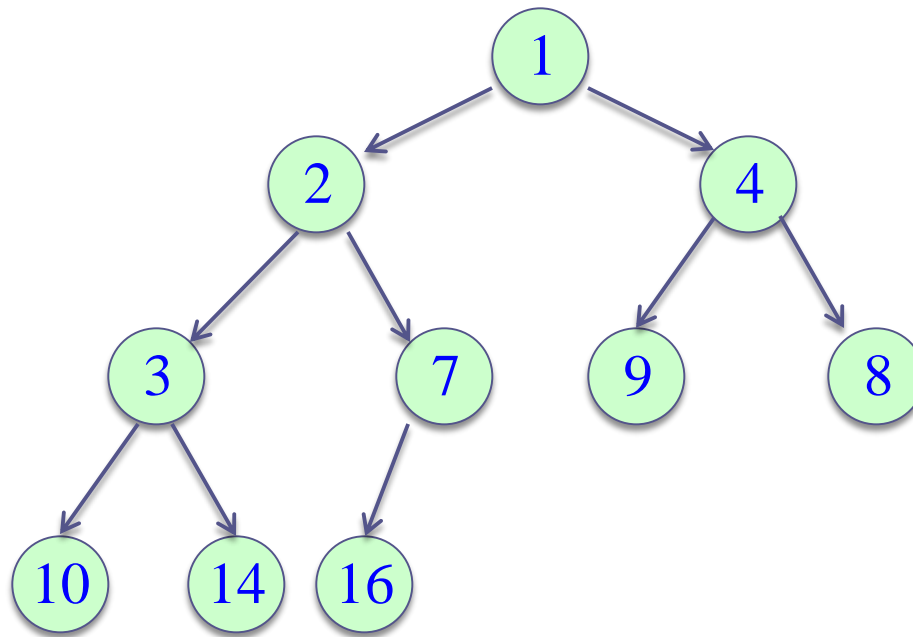
Heap Data Structure



Complete binary tree

- Completely filled on all levels except possibly the lowest level
- The lowest level is filled from left to right

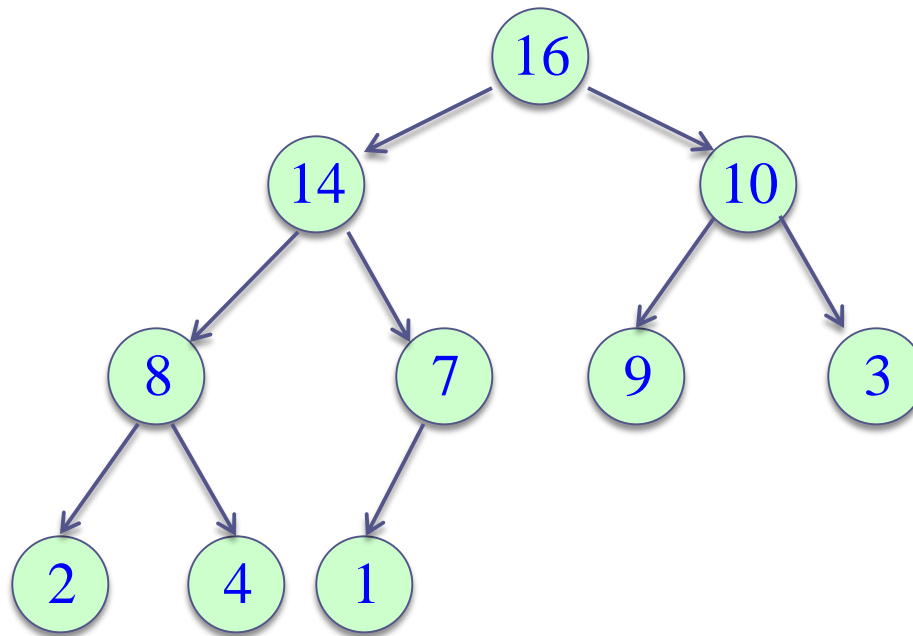
Heap Property: Min-Heap



The smallest element in any subtree is the root element in a min-heap

Min heap: For every node i other than root, $A[\text{parent}(i)] \leq A[i]$
→ Parent node is always smaller than the child nodes

Heap Property: Max-Heap



*The largest element
in any subtree is the root
element in a max-heap*

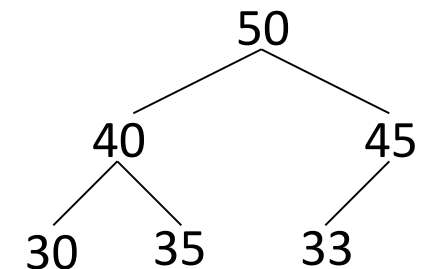
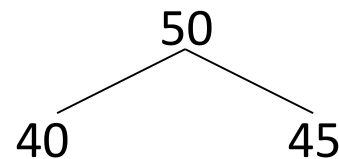
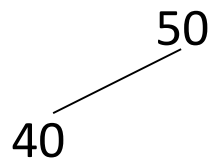
We will focus on max-heaps

Max heap: For every node i other than root, $A[\text{parent}(i)] \geq A[i]$
→ Parent node is always larger than the child nodes

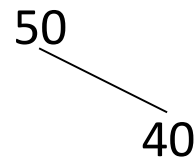
Differences between a Heap and a BST

- A heap is **NOT** a binary search tree.
 1. A BST can be seen as sorted, but a heap is ordered in much weaker sense.
 - Although it is not sorted, the order of a heap is sufficient for the efficient implementation of priority queue operations.
 2. A BST has different shapes, but **a heap is always complete binary tree.**

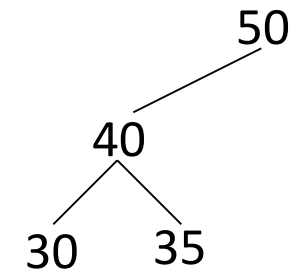
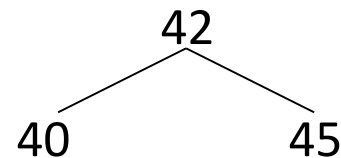
HEAPS



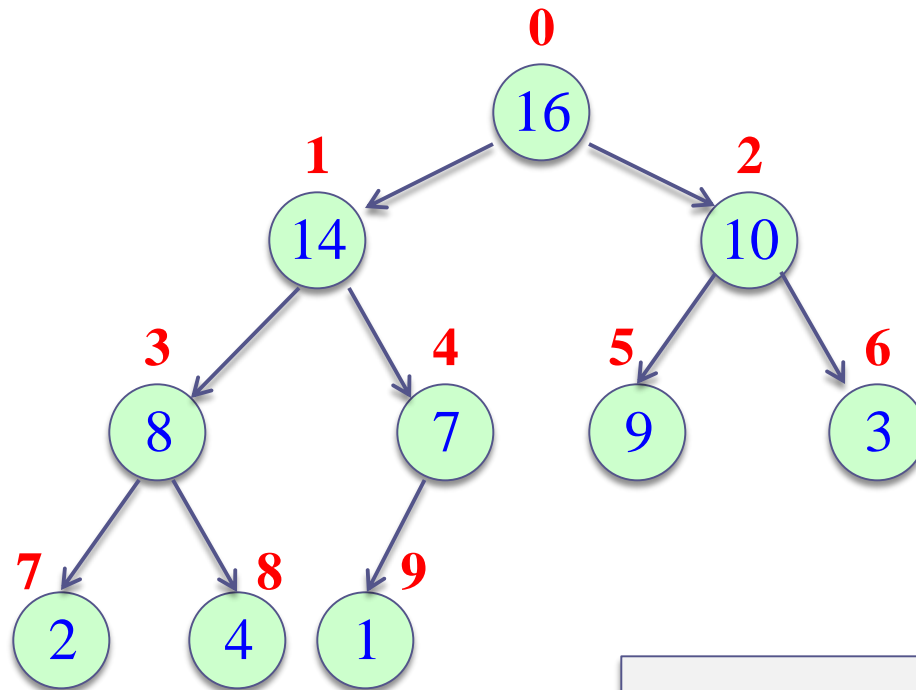
NOT HEAPS



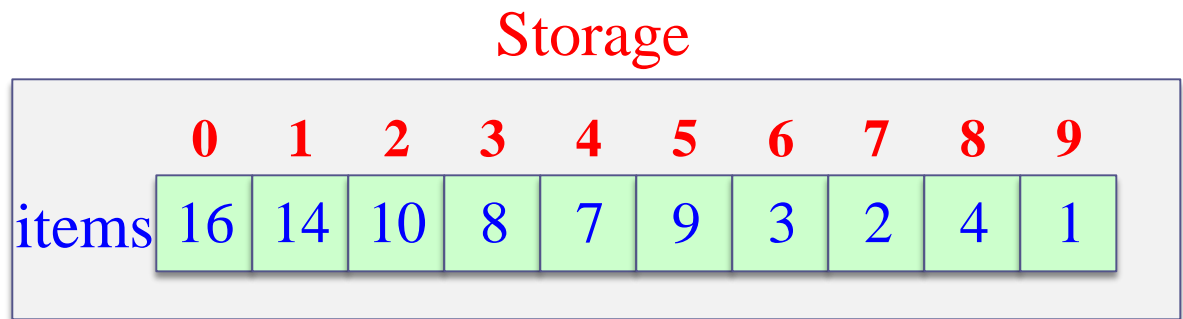
not complete binary tree



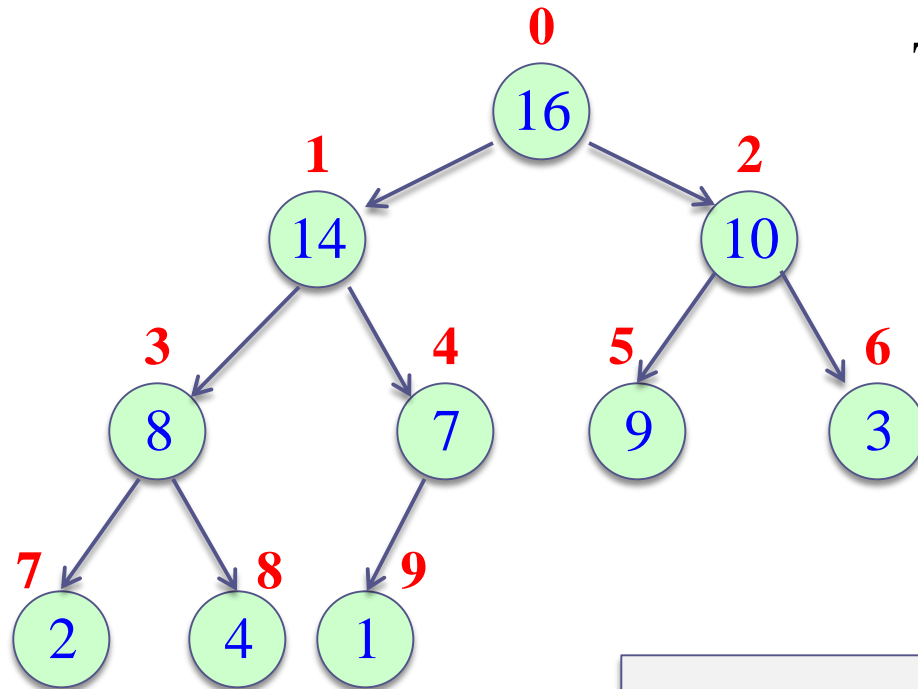
Heap Data Structure



Heap can be stored in a linear array



Heap Data Structure



The links in the heap are implicit:

$$\text{left}(i) = 2i + 1$$

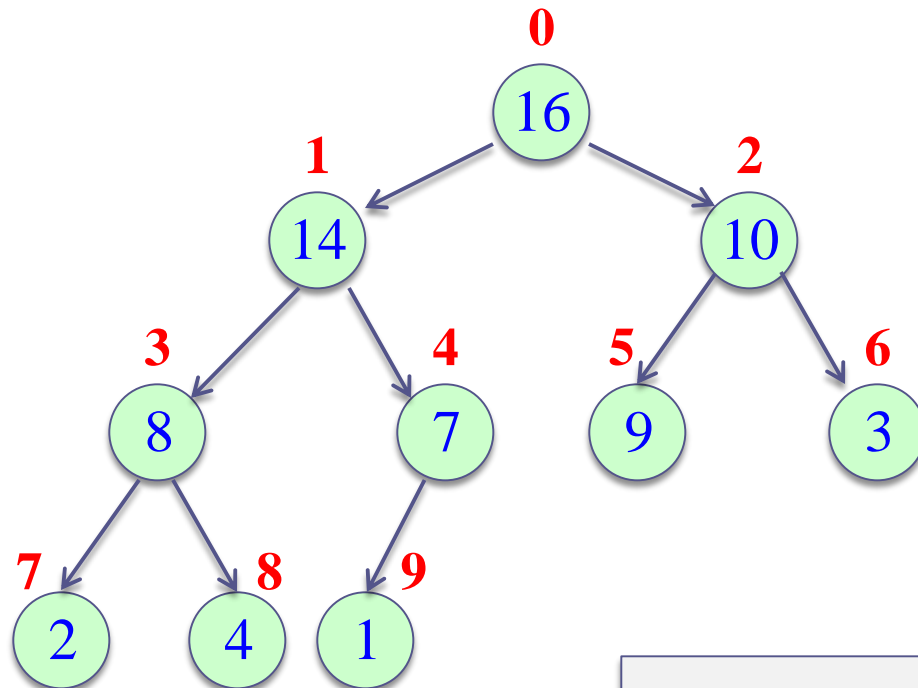
$$\text{right}(i) = 2i + 2$$

$$\text{parent}(i) = \lfloor (i - 1) / 2 \rfloor$$

Storage

	0	1	2	3	4	5	6	7	8	9
items	16	14	10	8	7	9	3	2	4	1

Heap Data Structure



$$\text{left}(i) = 2i + 1$$

e.g. Left child of node 3 has index 7

$$\text{right}(i) = 2i + 2$$

e.g. Right child of node 1 has index 4

$$\text{parent}(i) = \lfloor (i - 1) / 2 \rfloor$$

e.g. Parent of node 6 has index 2

	0	1	2	3	4	5	6	7	8	9
items	16	14	10	8	7	9	3	2	4	1

Heap Data Structure

- `items[0]` is always the root element
- Array `items` has two attributes:
 - ▣ `MAX_SIZE`: Size of the memory allocated for array `items`
 - ▣ `size`: The number elements in heap at a given time

$$\text{size} \leq \text{MAX_SIZE}$$

Major Heap Operations

- Two major heap operations are insertion and deletion.

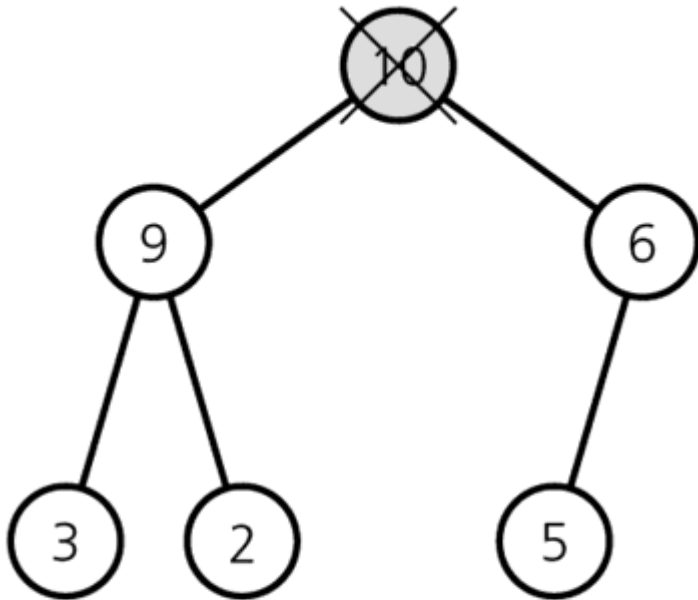
Insertion

- Inserts a new item into a heap.
- After the insertion, the heap must satisfy the heap properties.

Deletion

- Retrieves and deletes the root of the heap.
- After the deletion, the heap must satisfy the heap properties.

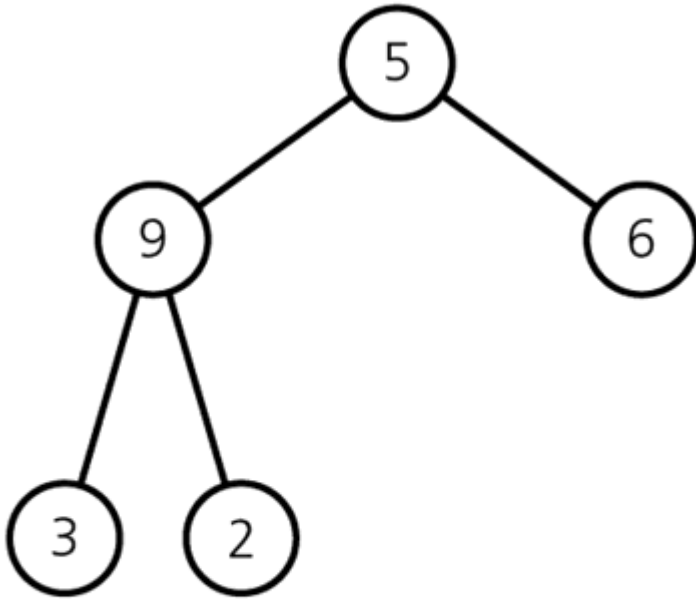
Heap Delete – First Step



- The first step of **heapDelete** is to retrieve and delete the root.
 - This creates two disjoint heaps.

0	10
1	9
2	6
3	3
4	2
5	5

Heap Delete – Second Step

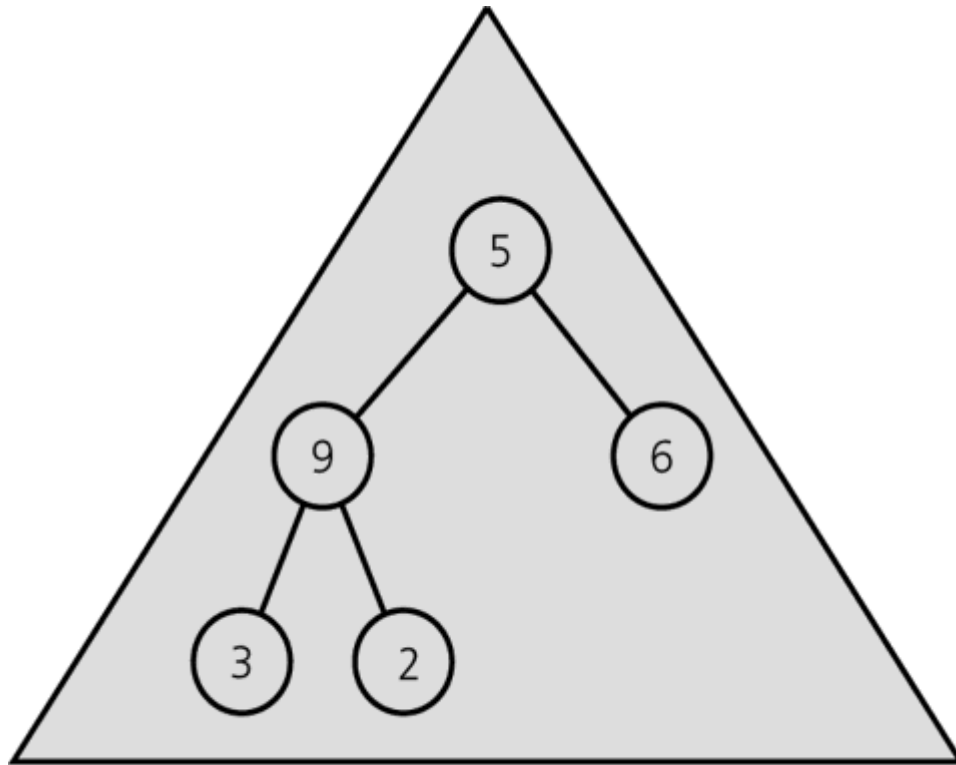


- Move the last item into the root.
- The resulting structure may not be heap; it is called as **semiheap**.

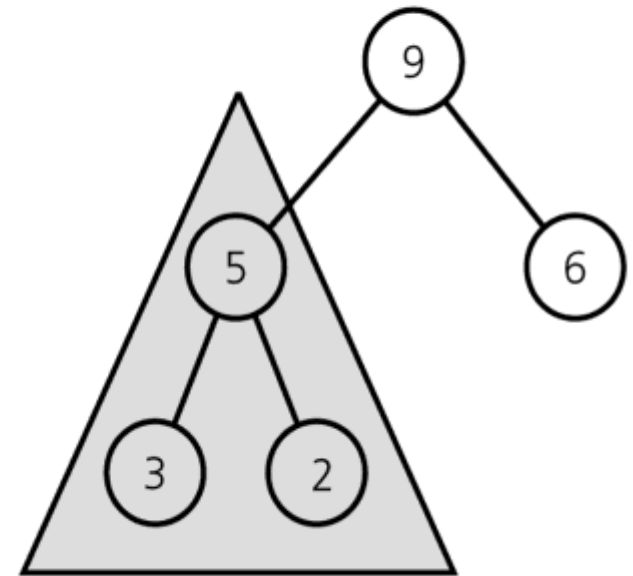
0	5
1	9
2	6
3	3
4	2

Heap Delete – Last Step

The last step of *heapDelete* transforms the semiheap into a heap.



First semiheap passed to *heapRebuild*



Second semiheap passed to *heapRebuild*

Recursive calls to **heapRebuild**

Heap Delete

heapDelete (items, size)

$\text{max} \leftarrow \text{items}[0]$

$\text{items}[0] \leftarrow \text{items}[\text{size}-1]$

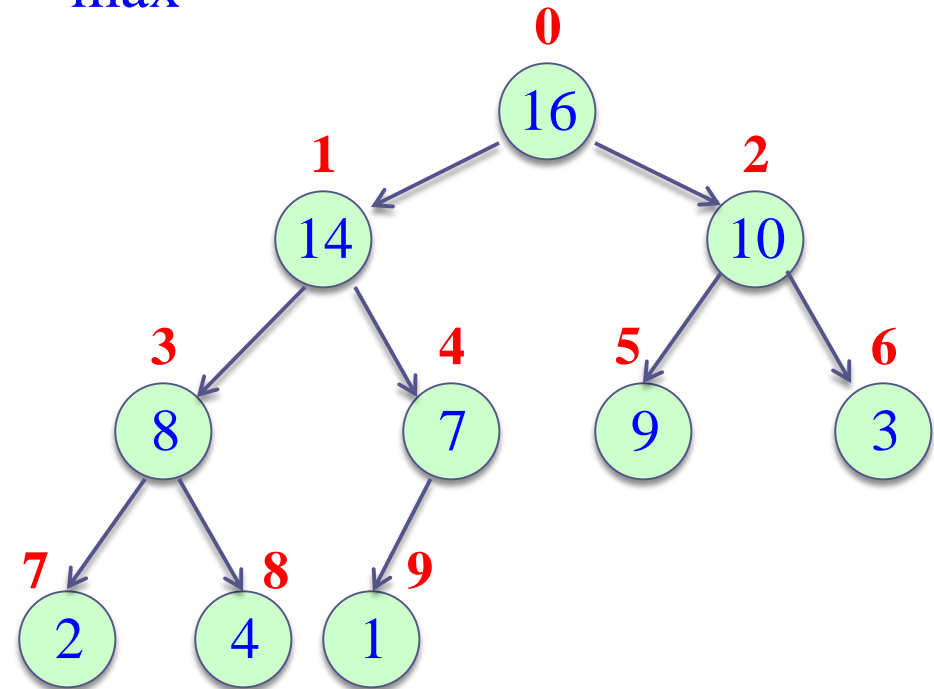
$\text{size} \leftarrow \text{size} - 1$

heapRebuild(items, 0, size)

return max

*Return the max element,
and reorganize the heap
to maintain heap property*

max=



only decrement size in the array
does not actually delete the node

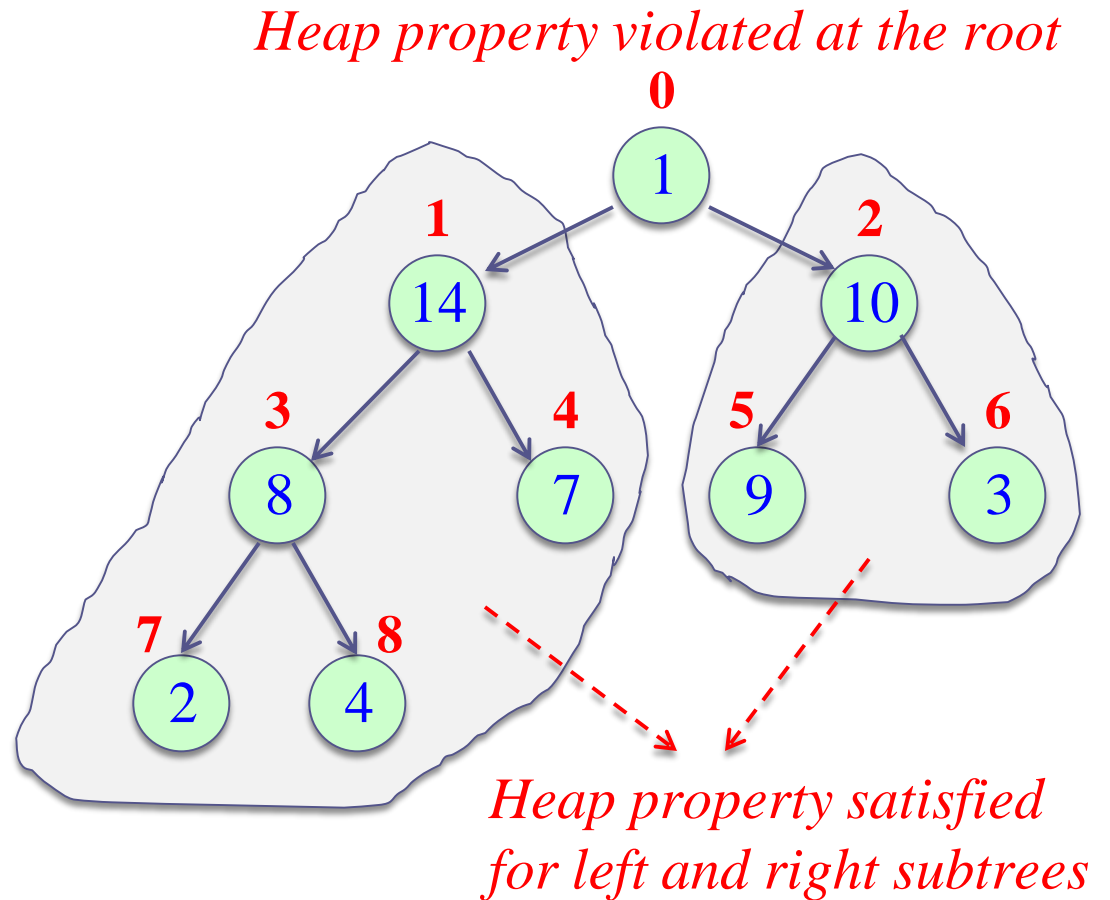
Rebuild Heap

Maintaining heap property:

Subtrees rooted at $\text{left}(i)$ and $\text{right}(i)$ are already heaps.

But, $\text{items}[i]$ may violate the heap property (i.e., may be smaller than its children)

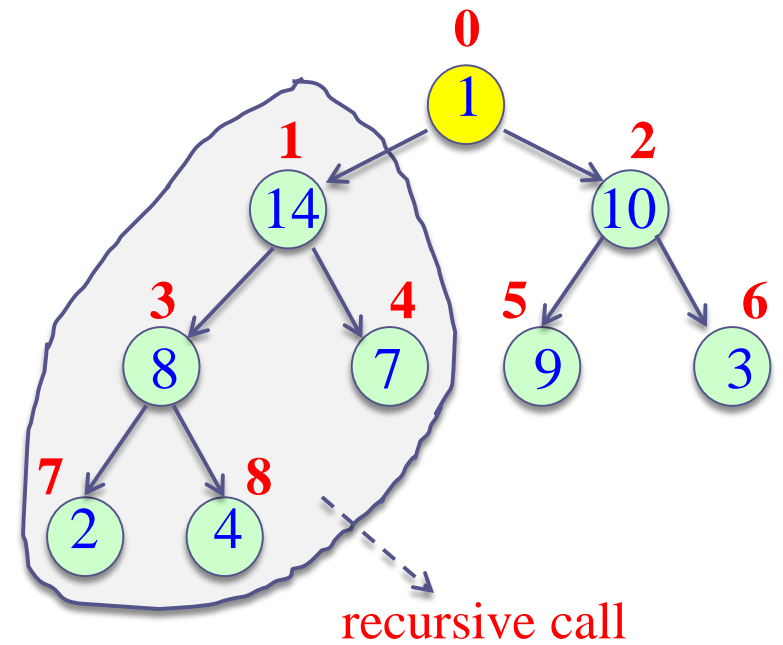
Idea: Float down the value at $\text{items}[i]$ in the heap so that subtree rooted at i becomes a heap.



Rebuild Heap

select larger child

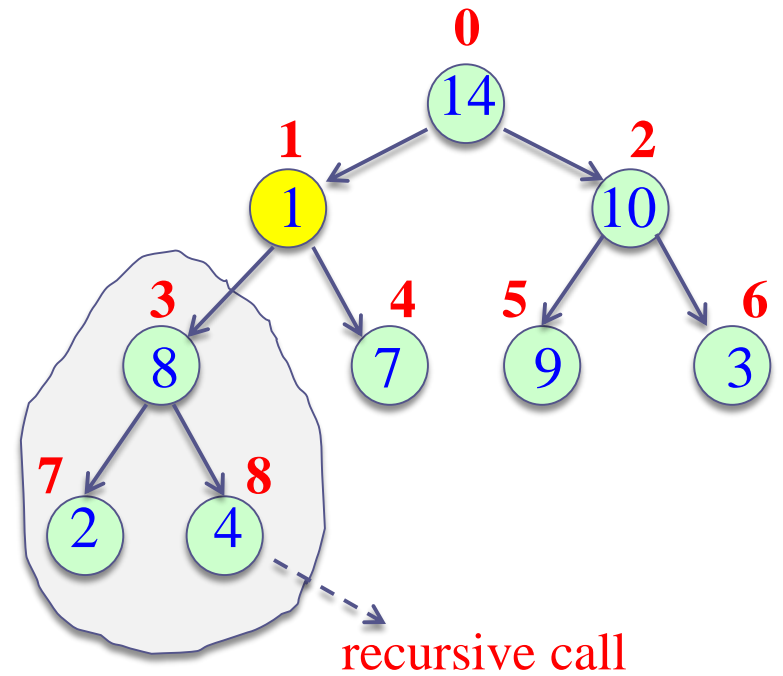
rebuildHeap(items, 0, 9)



Rebuild Heap

recursive call:

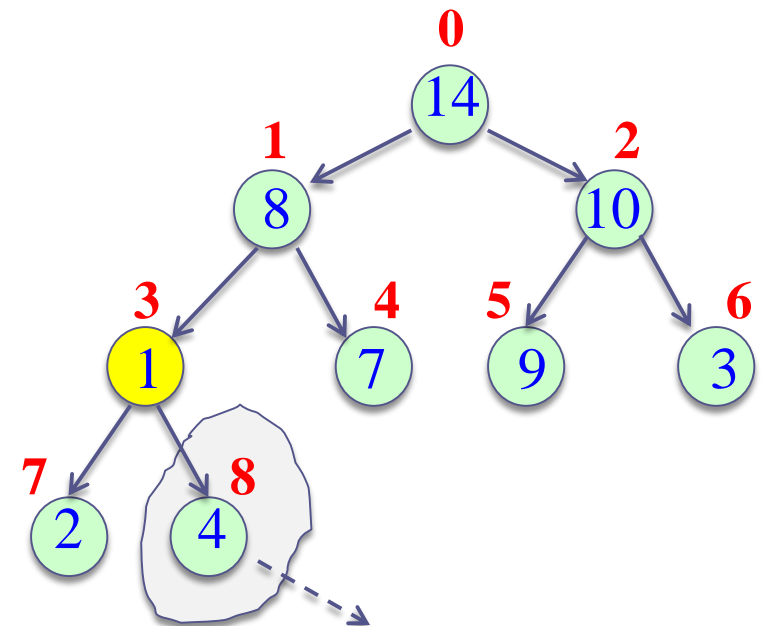
rebuildHeap(items, 1, 9)



Rebuild Heap

recursive call:

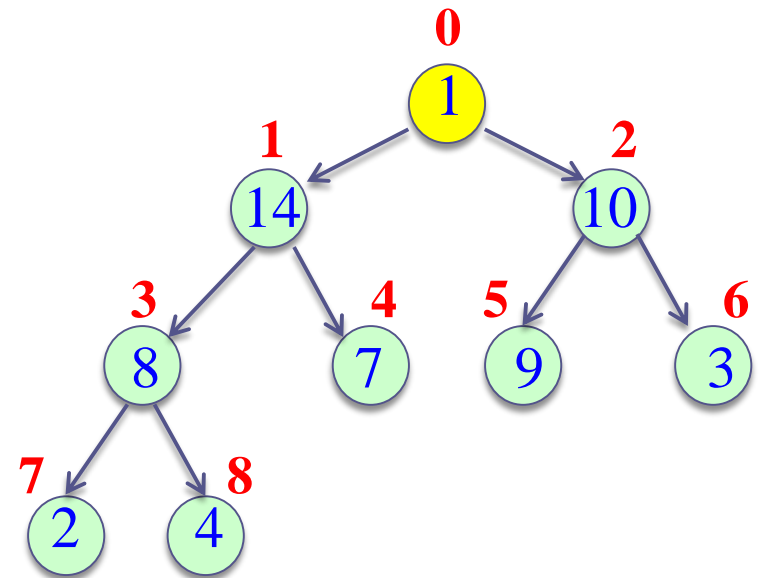
rebuildHeap(items, 3, 9)



recursive call
(base case)

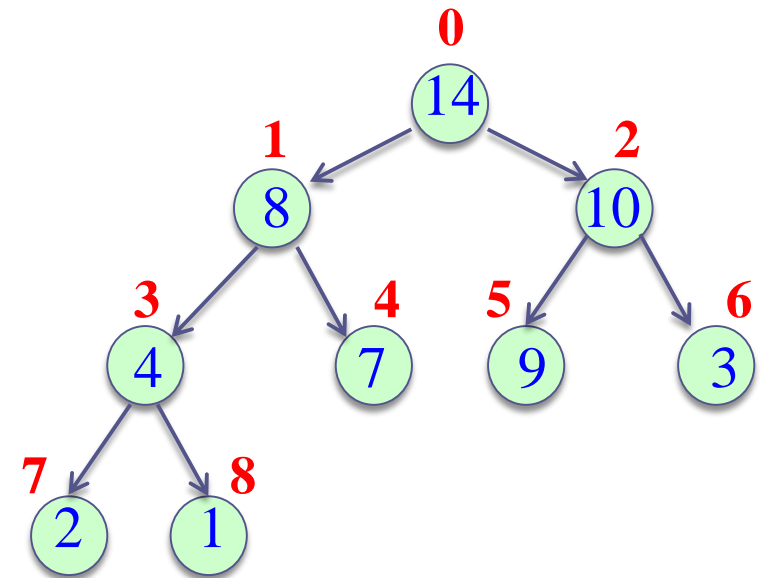
Rebuild Heap: Summary (Floating Down the Value)

rebuildHeap(items, 0, 9)

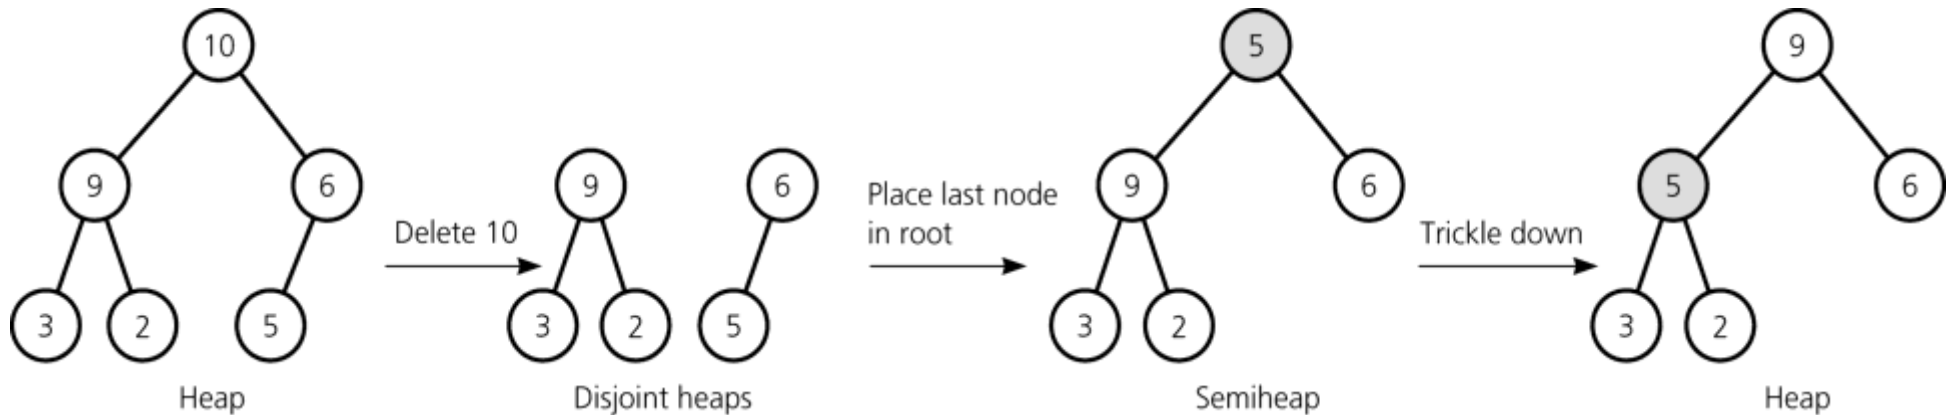


Heap Operations: Rebuild Heap

after rebuildHeap:



Heap Delete

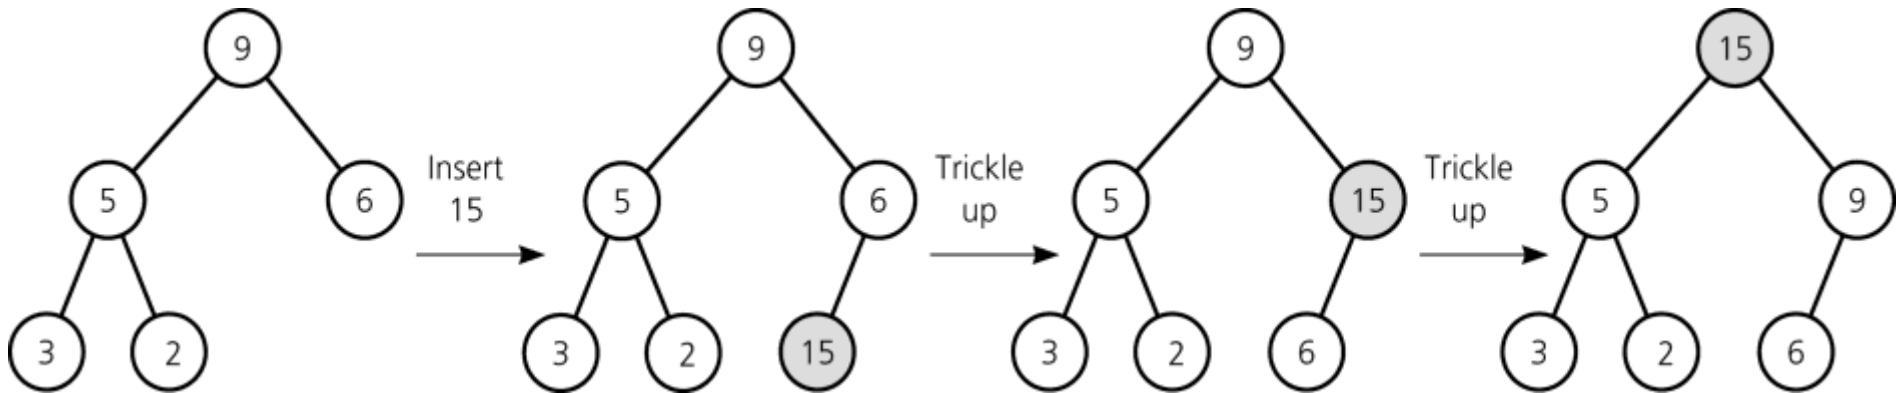


ANALYSIS

- Since the height of a complete binary tree with n nodes is always $\lceil \log_2(n+1) \rceil$
 - ➔ **heapDelete** is **$O(\log_2 n)$** worst case and average case

Heap Insert

A new item is inserted at the bottom of the tree, and it trickles up to its proper place



ANALYSIS

- Since the height of a complete binary tree with n nodes is always $\lceil \log_2(n+1) \rceil$

➔ **heapInsert** is **$O(\log_2 n)$**

worst case and average case

Heap Implementation

```
constint MAX_HEAP = maximum-size-of-heap;
#include "KeyedItem.h" // definition of KeyedItem
typedef KeyedItem HeapItemType;

class Heap {
public:
    Heap(); // default constructor
    // copy constructor and destructor are supplied by the compiler

    bool heapIsEmpty() const;
    void heapInsert(const HeapItemType& newItem) throw(HeapException);
    void heapDelete(HeapItemType& rootItem) throw(HeapException);

protected:
    void heapRebuild(int root); // Converts the semiheap rooted at
                                // index root into a heap

private:
    HeapItemType items[MAX_HEAP]; // array of heap items
    int size; // number of heap items
};
```

Heap Implementation

```
// Default constructor
Heap::Heap() : size(0) {

}

bool Heap::isEmpty() const {
    return (size == 0);
}
```

Heap Implementation -- *heapInsert*

```
void Heap::heapInsert(const HeapItemType&newItem) throw(HeapException) {  
  
    if (size >= MAX_HEAP)  
        throw HeapException("HeapException: Heap full");  
  
    // Place the new item at the end of the heap  
    items[size] = newItem;  
  
    // Trickle new item up to its proper position  
    int place = size;  
    int parent = (place - 1)/2;    integer division, root index no matter left or right child  
    while ( (place > 0) && (items[place].getKey() > items[parent].getKey()) ) {  
        HeapItemType temp = items[parent];  
        items[parent] = items[place];  
        items[place] = temp;  
  
        place = parent;  
        parent = (place - 1)/2;  
    }  
    ++size;  
}
```

can be recursive
heapRebuild(lastIndex)
modify heapRebuild to trickle up
instead of root to leaf

Heap Implementation -- *heapDelete*

```
Void Heap::heapDelete(HeapItemType&rootItem) throw(HeapException) {  
    if (heapIsEmpty())  
        throwHeapException("HeapException: Heap empty");  
    else {  
        rootItem = items[0];  
        items[0] = items[--size];  
        heapRebuild(0);  
    }  
}
```

Heap Implementation -- *heapRebuild*

```
void Heap::heapRebuild(int root) {
    int child = 2 * root + 1;    // index of root's left child, if any
    if ( child < size ) {
        // root is not a leaf so that it has a left child
        int rightChild = child + 1;    // index of a right child, if any

        // If root has right child, find larger child
        if ( (rightChild < size) &&
            (items[rightChild].getKey() > items[child].getKey()) )
            child = rightChild;    // index of larger child

        // If root's item is smaller than larger child, swap values
        if ( items[root].getKey() < items[child].getKey() ) {
            HeapItemType temp = items[root];
            items[root] = items[child];
            items[child] = temp;

            // transform the new subtree into a heap
            heapRebuild(child);
        }
    }
}
```

Heap Implementation of *PriorityQueue*

- The heap implementation of the priority queue is straightforward
 - Since the heap operations and the priority queue operations are the same.
- When we use the heap,
 - Insertion and deletion operations of the priority queue will be $O(\log_2 n)$.

Heap Implementation of *PriorityQueue*

```
#include "Heap.h"// ADT heap operations
typedef HeapItemType PQItemType;

class PriorityQueue {
public:
    // default constructor, copy constructor, and destructor
    // are supplied by the compiler

    // priority-queue operations:
    bool pqIsEmpty() const;
    void pqInsert(const PQItemType& newItem) throw (PQException);
    void pqDelete(PQItemType& priorityItem) throw (PQException);

private:
    Heap h;
};
```

Heap Implementation of *PriorityQueue*

```
bool PriorityQueue::pqIsEmpty() const {
    return h.heapIsEmpty();
}

void PriorityQueue::pqInsert(const PQItemType& newItem) throw (PQException) {
    try {
        h.heapInsert(newItem);
    }
    catch (HeapException e) {
        throw PQException("Priority queue is full");
    }
}

void PriorityQueue::pqDelete(PQItemType& priorityItem) throw (PQException) {
    try {
        h.heapDelete(priorityItem);
    }
    catch (HeapException e) {
        throw PQException("Priority queue is empty");
    }
}
```


Heap or Binary Search Tree?

Heapsort

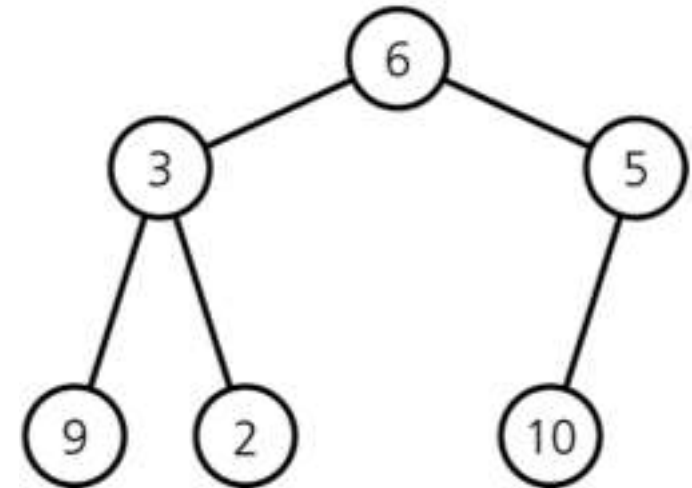
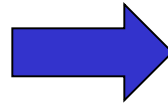
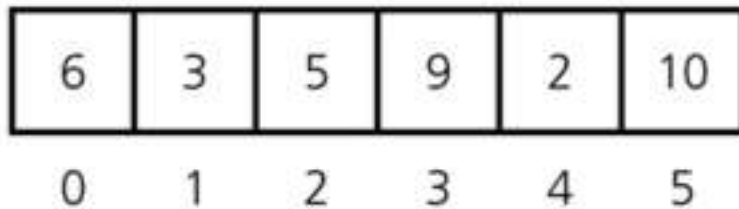
We can make use of a heap to sort an array:

1. Create a heap from the given initial array with n items.
2. Swap the root of the heap with the last element in the heap.
3. Now, we have a semiheap with $n-1$ items, and a sorted array with one item.
4. Using `heapRebuild` convert this semiheap into a heap. Now we will have a heap with $n-1$ items.
5. Repeat the steps 2-4 as long as the number of items in the heap is more than 1.

Heapsort -- Building a heap from an array

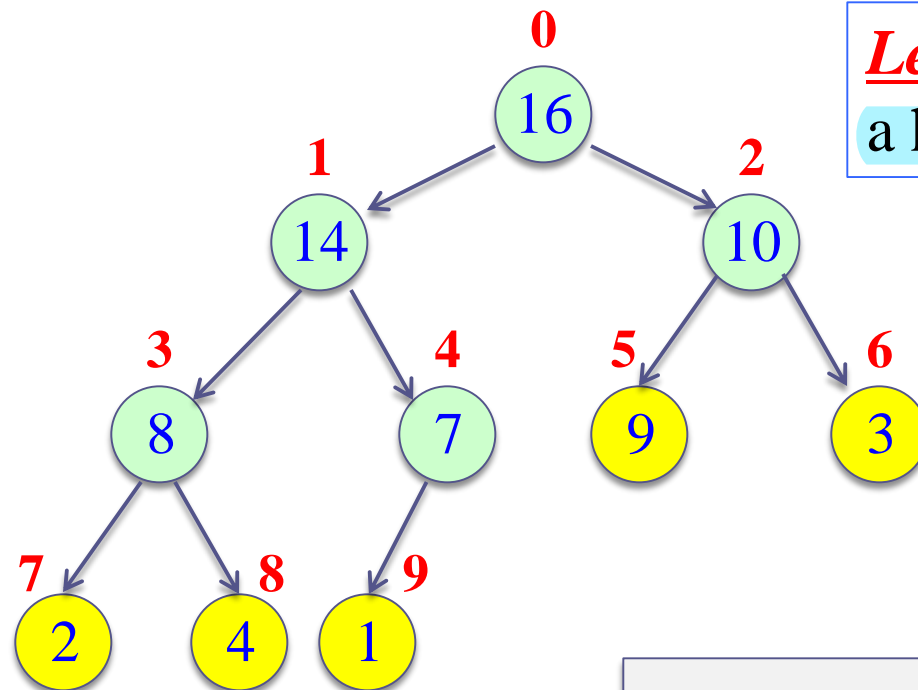
A heap corresponding to *anArray*

The initial contents of *anArray*



```
for (index = n - 1 ; index >= 0 ; index--) {  
    // Invariant: the tree rooted at index is a semiheap  
    heapRebuild(anArray, index, n)  
    // Assertion: the tree rooted at index is a heap.  
}
```

Where are the leaves stored?



Lemma: The last $\lceil n/2 \rceil$ nodes of a heap are *all leaves*

it is pointless to call heapRebuild after the index 4

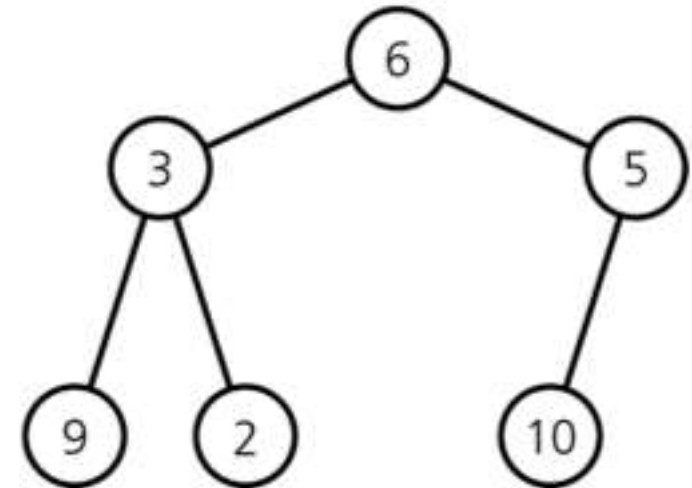
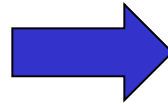
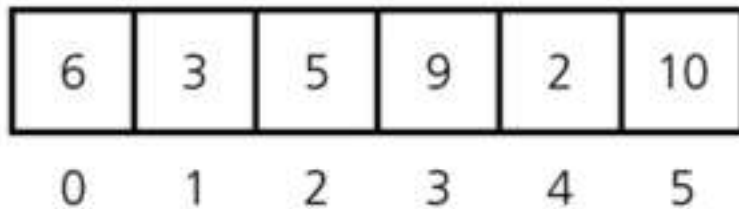
Storage

	0	1	2	3	4	5	6	7	8	9
A	16	14	10	8	7	9	3	2	4	1

Heapsort -- Building a heap from an array

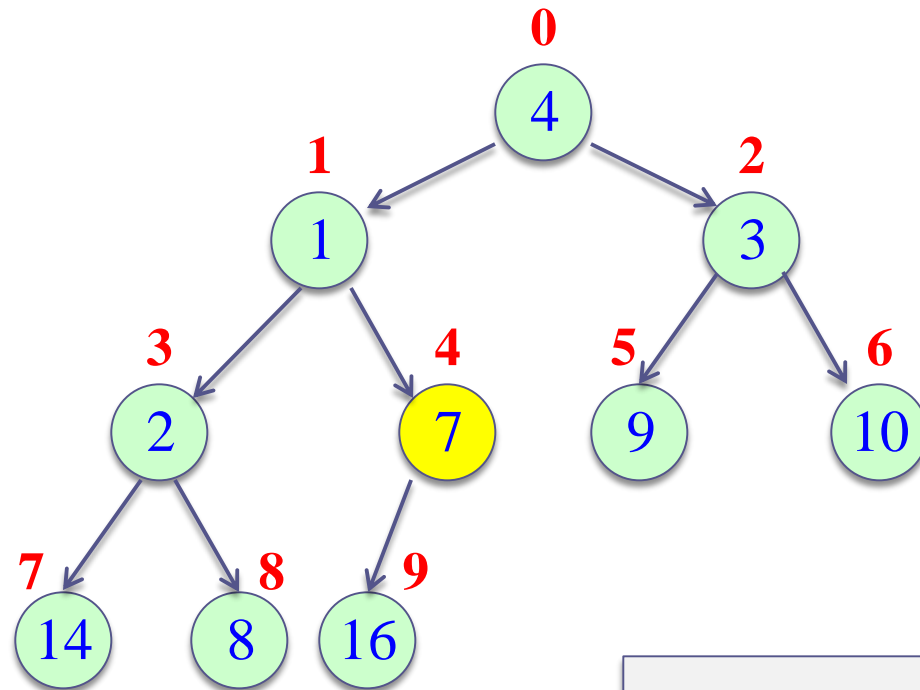
A heap corresponding to *anArray*

The initial contents of *anArray*



```
for (index = (n/2) - 1 ; index >= 0 ; index--) {  
    → MORE EFFICIENT  
    // Invariant: the tree rooted at index is a semiheap  
    heapRebuild(anArray, index, n)  
    // Assertion: the tree rooted at index is a heap.  
}
```

Build-Heap: Example

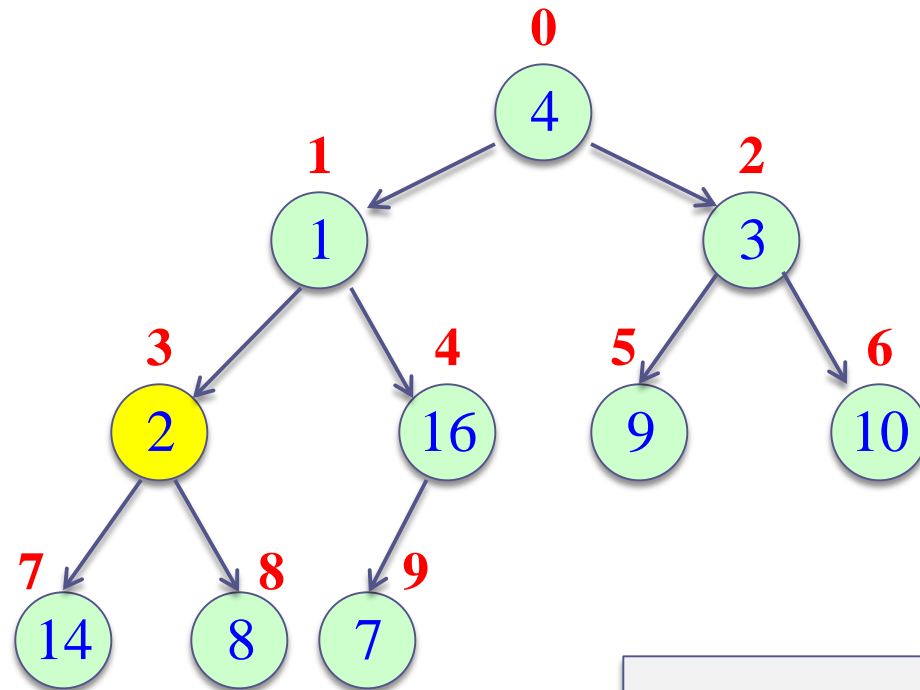


index=4

heapRebuild(A, 4, 10)

	0	1	2	3	4	5	6	7	8	9
A	4	1	3	2	7	9	10	14	8	16

Build-Heap: Example

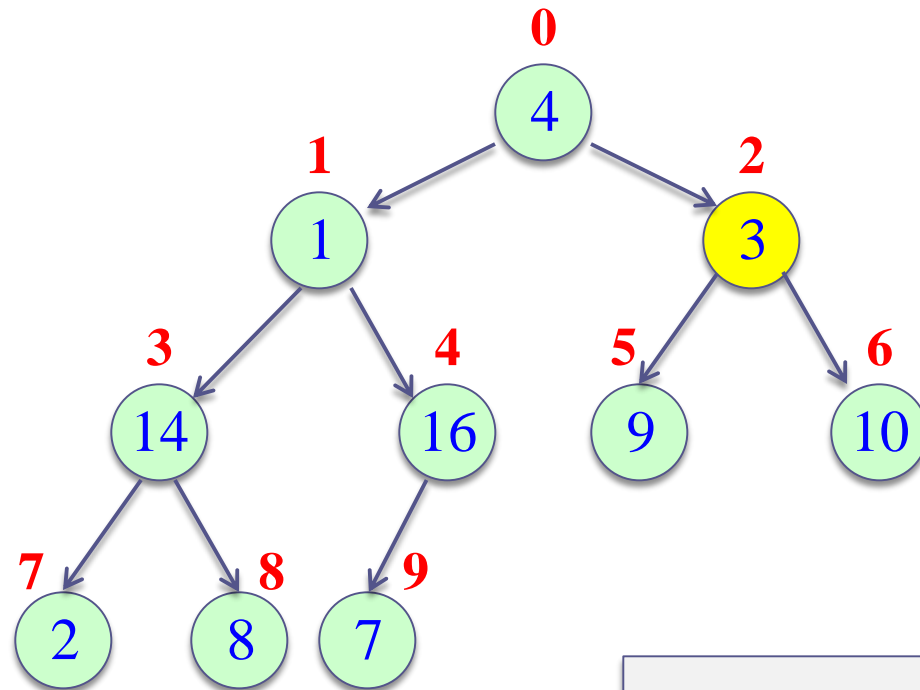


index=3

heapRebuild(A, 3, 10)

	0	1	2	3	4	5	6	7	8	9
A	4	1	3	2	16	9	10	14	8	7

Build-Heap: Example

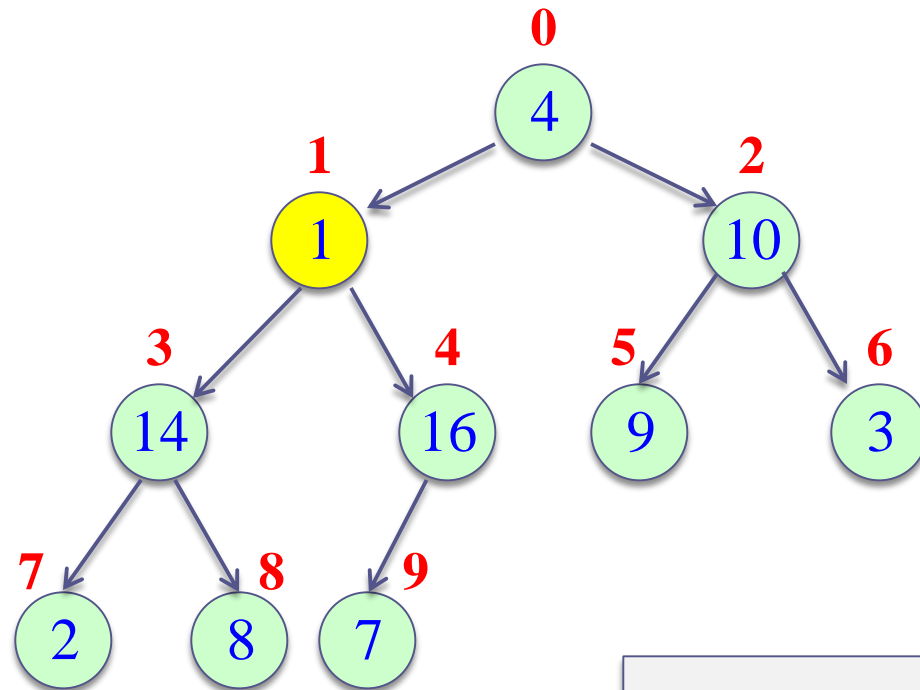


index=2

heapRebuild(A, 2, 10)

	0	1	2	3	4	5	6	7	8	9
A	4	1	3	14	16	9	10	2	8	7

Build-Heap: Example

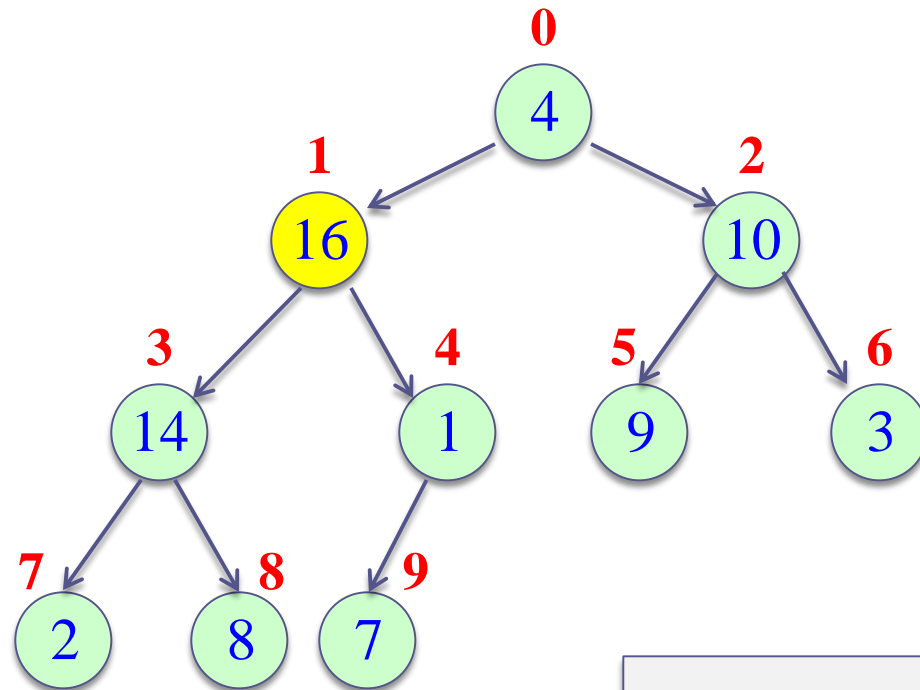


index=1

heapRebuild(A, 1, 10)

	0	1	2	3	4	5	6	7	8	9
A	4	1	10	14	16	9	3	2	8	7

Build-Heap: Example

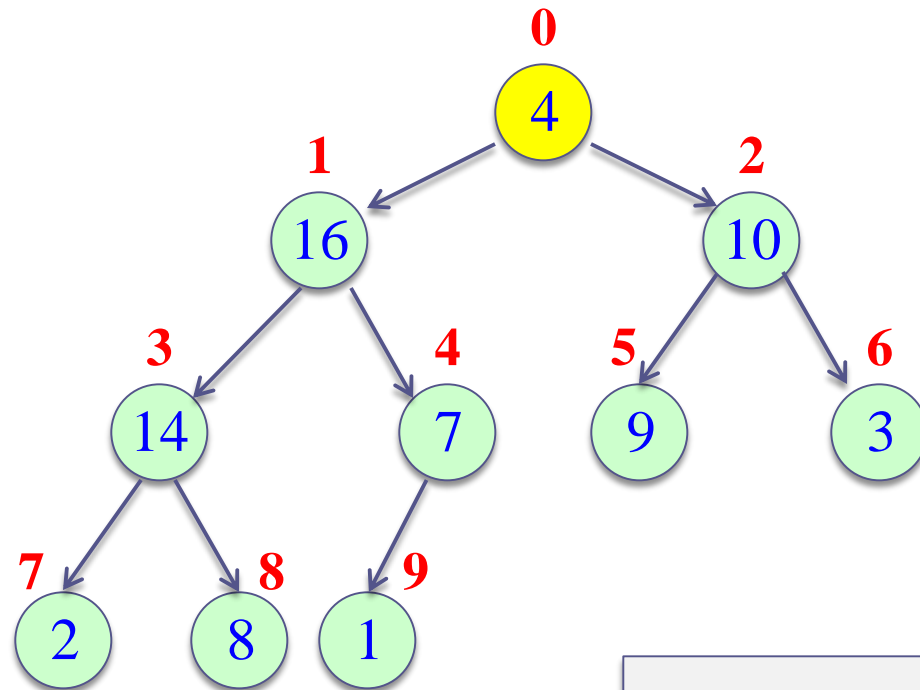


index=1 (cont'd)

heapRebuild(A, 1, 10)

	0	1	2	3	4	5	6	7	8	9
A	4	16	10	14	1	9	3	2	8	7

Build-Heap: Example

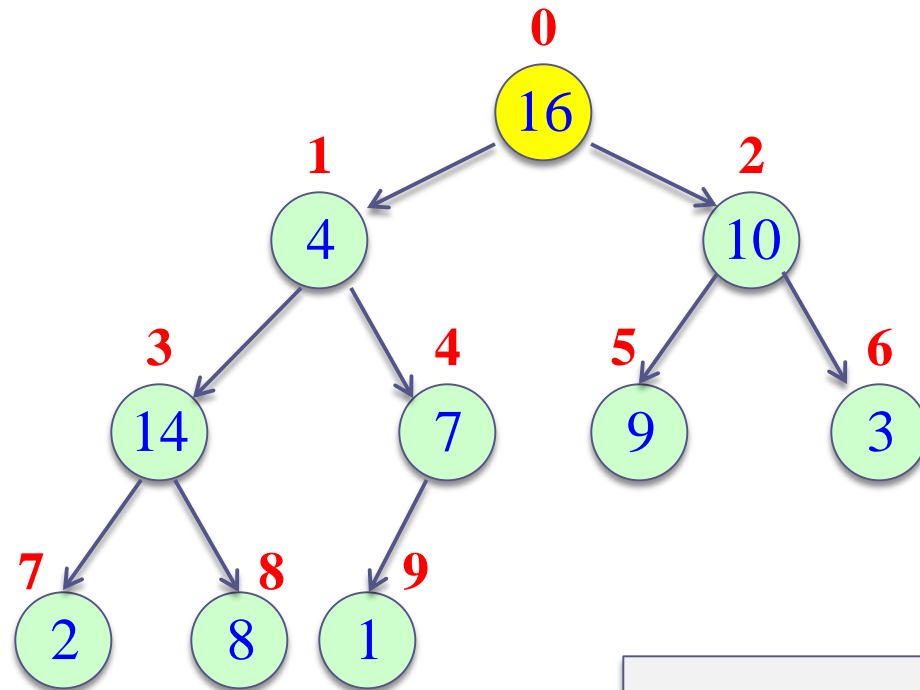


index=0

heapRebuild(A, 0, 10)

	0	1	2	3	4	5	6	7	8	9
A	4	16	10	14	7	9	3	2	8	1

Build-Heap: Example

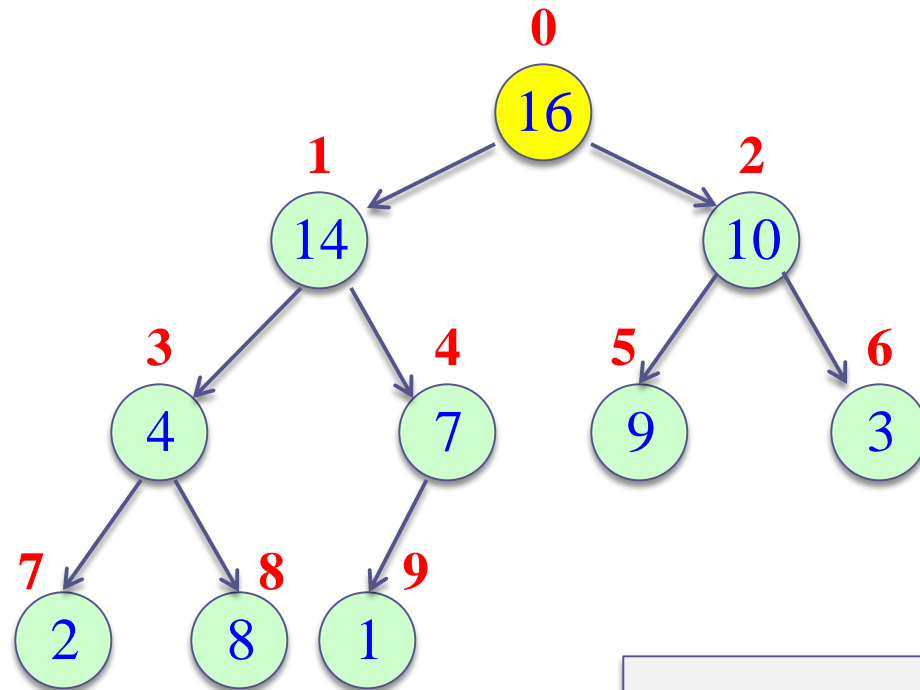


index=0 (cont'd)

heapRebuild(A, 0, 10)

	0	1	2	3	4	5	6	7	8	9
A	16	4	10	14	7	9	3	2	8	1

Build-Heap: Example



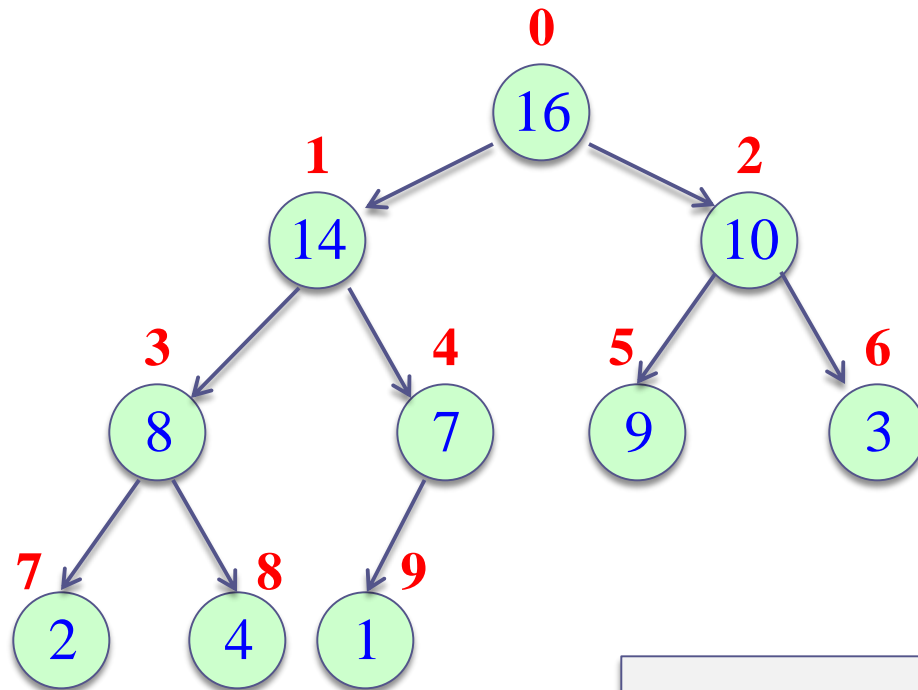
index=0 (cont'd)

heapRebuild(A, 0, 10)

	0	1	2	3	4	5	6	7	8	9
A	16	14	10	4	7	9	3	2	8	1

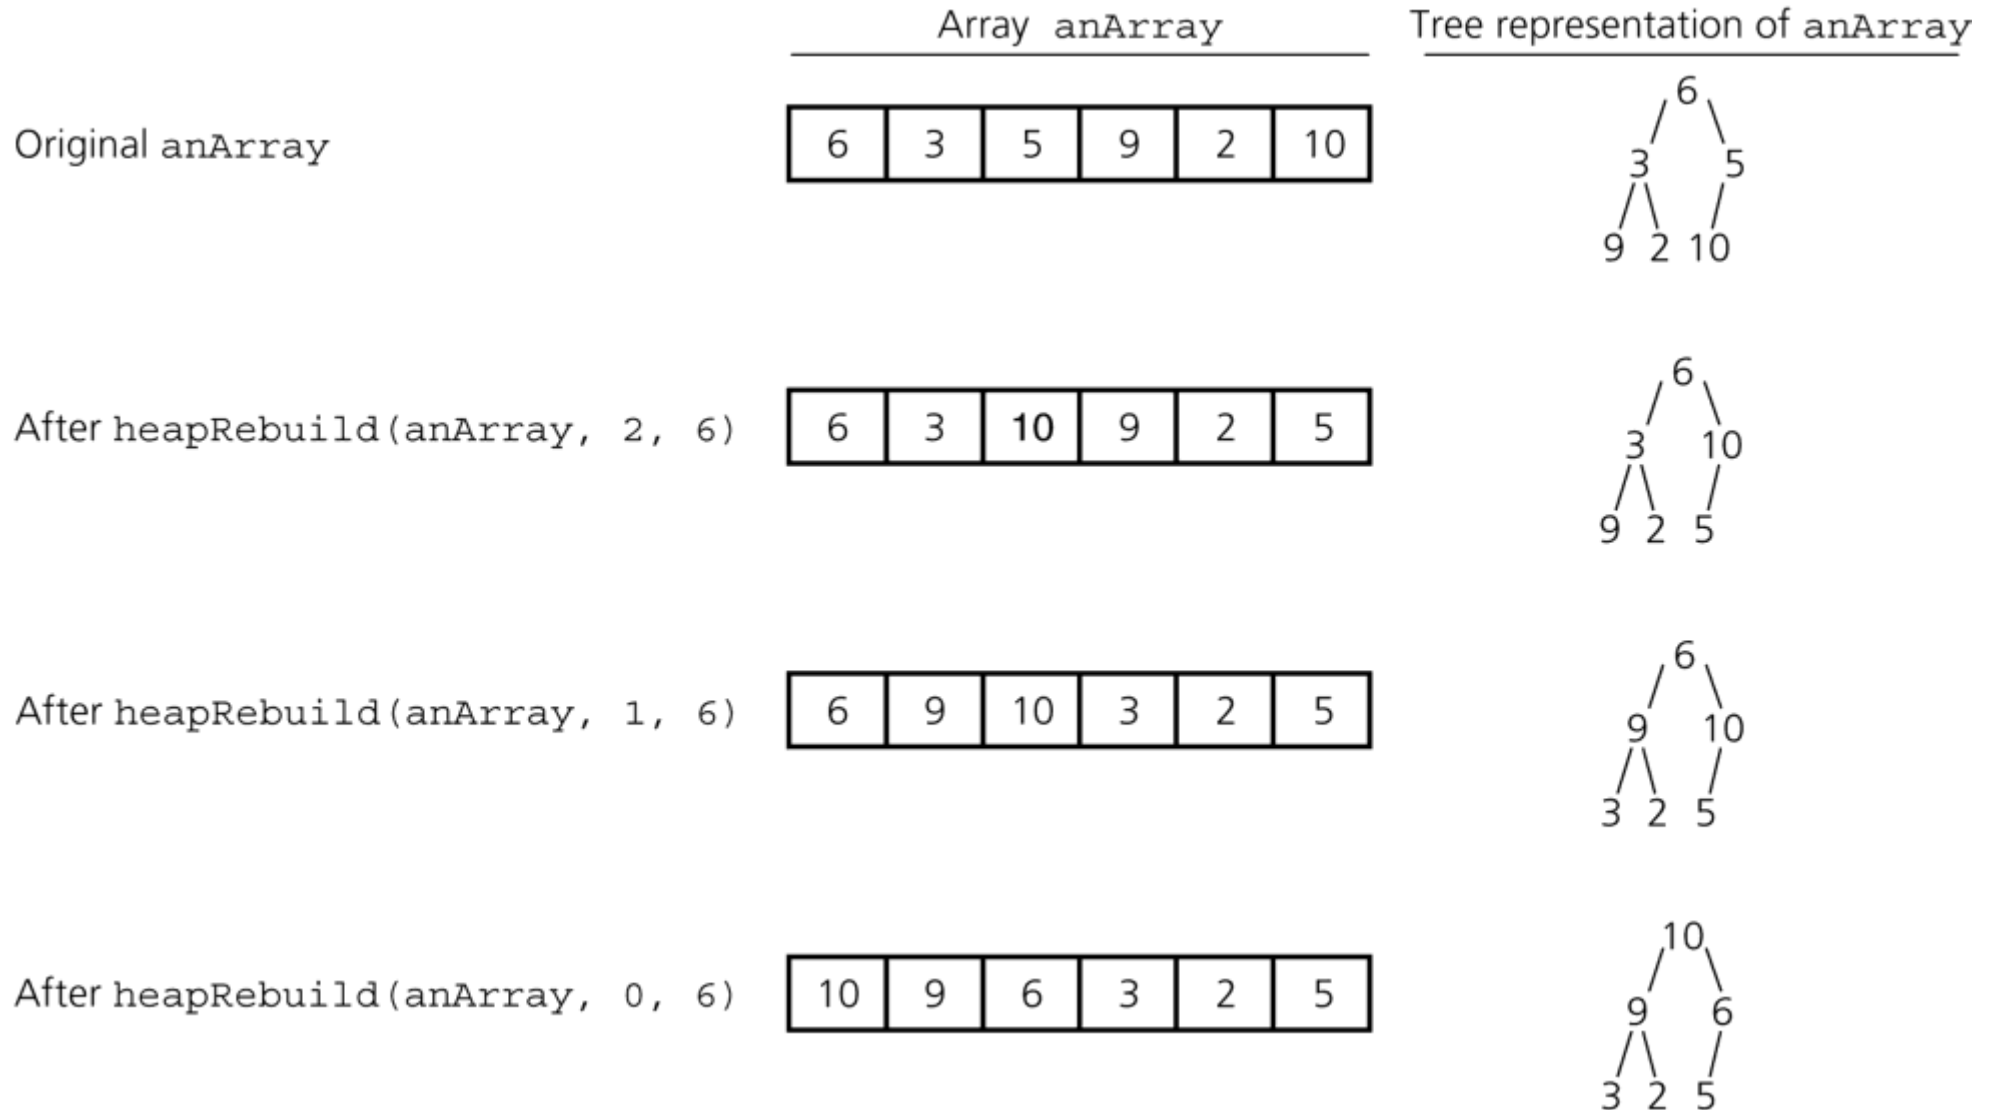
Build-Heap: Example

After Build-Heap



	0	1	2	3	4	5	6	7	8	9
A	16	14	10	8	7	9	3	2	4	1

Heapsort -- Building a heap from an array

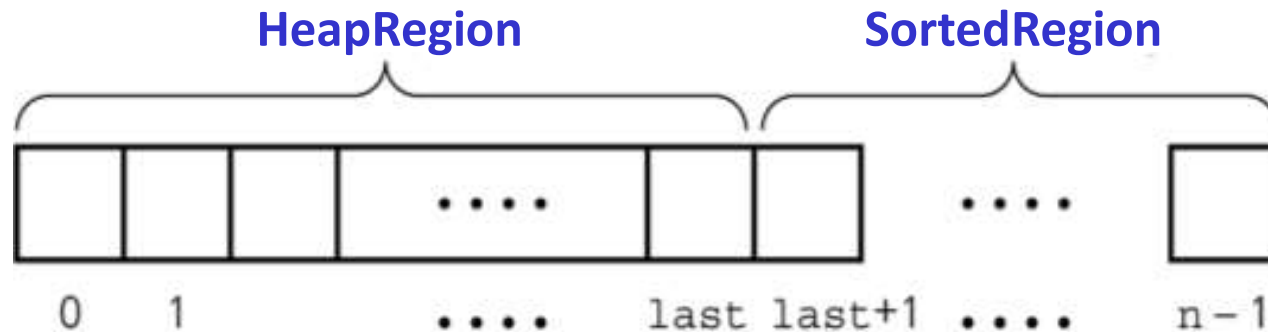


Heapsort

```
heapSort(inout anArray:ArrayType, in n:integer) {  
  
    // build an initial heap  
    for (index = (n/2) - 1 ; index >= 0 ; index--) O(n)  
        heapRebuild(anArray, index, n)  
  
    for (last = n-1 ; last > 0 ; last--) {  
        // invariant: anArray[0..last] is a heap,  
        // anArray[last+1..n-1] is sorted and  
        // contains the largest items of anArray.  
  
        swap anArray[0] and anArray[last] O(nlogn)  
  
        // make the heap region a heap again  
        heapRebuild(anArray, 0, last)  
    }  
}
```

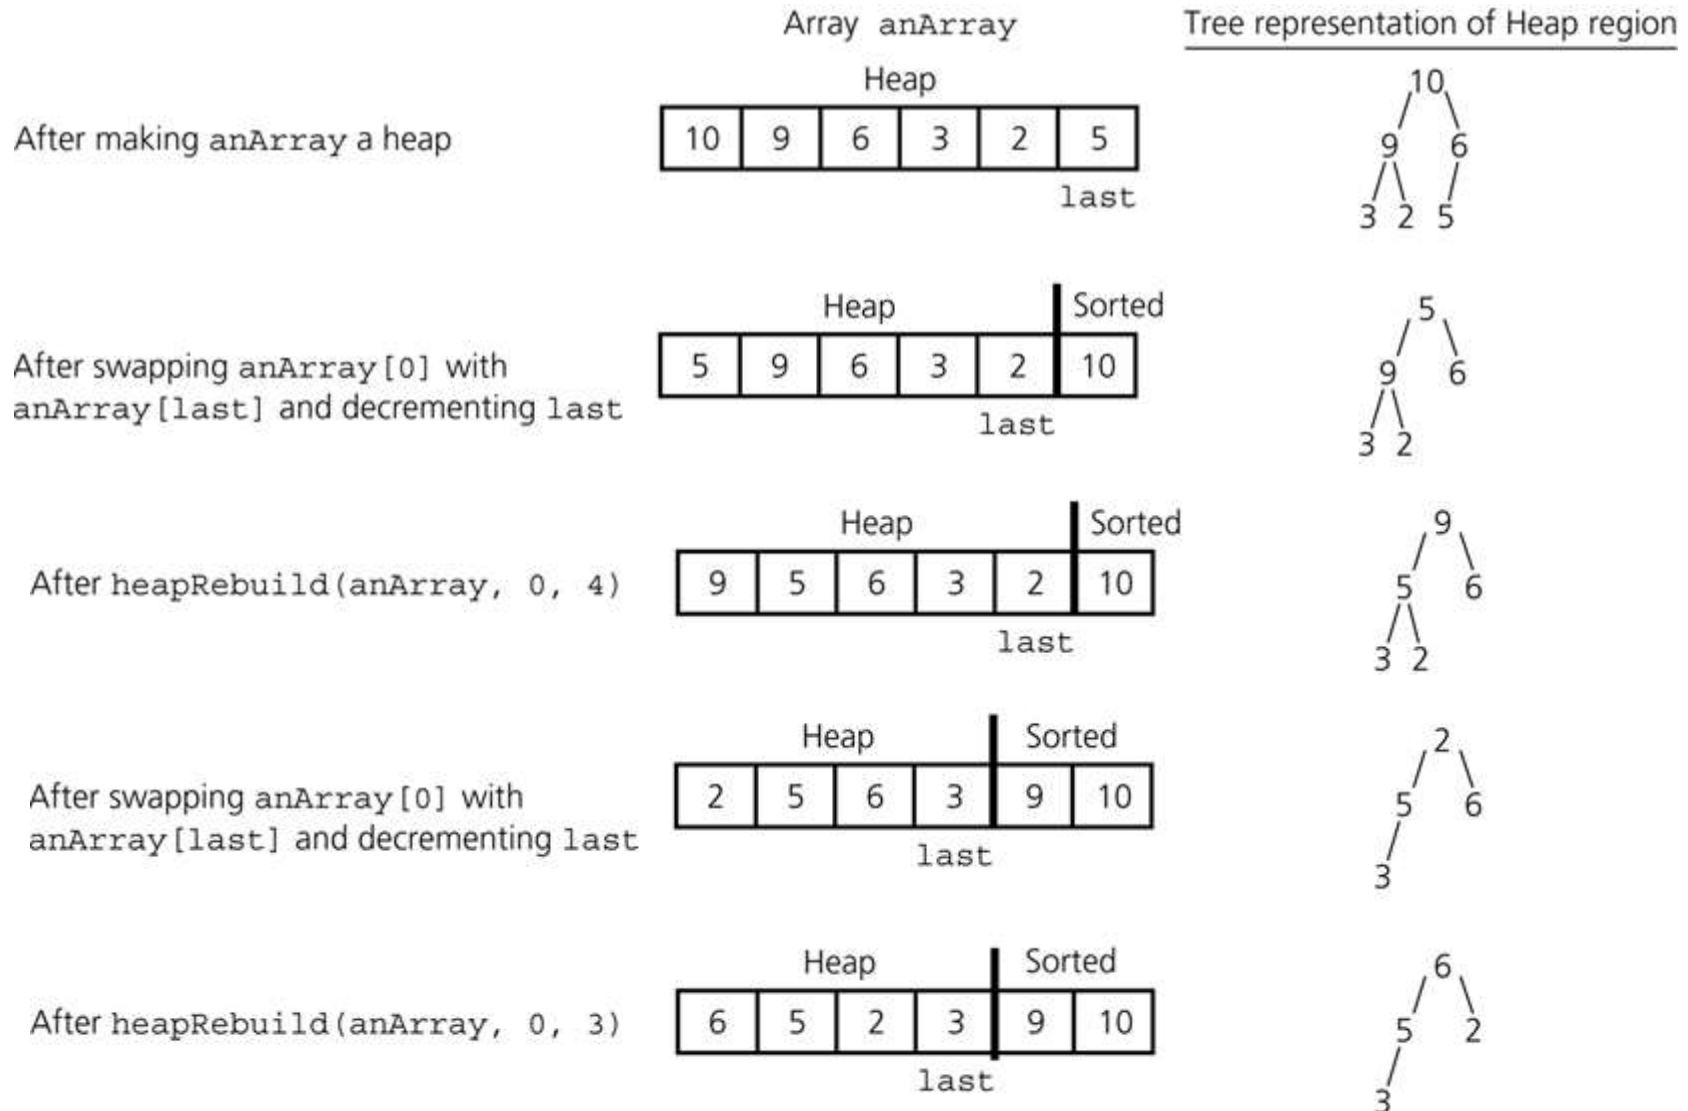

Heapsort

- Heapsort partitions an array into two regions.



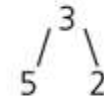
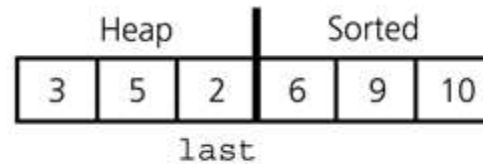
- Each step moves an item from the *HeapRegion* to *SortedRegion*.
- The invariant of the heapsort algorithm is:
 - After the *k*th step,
 - The *SortedRegion* contains the *k* largest value and they are in sorted order.
 - The items in the *HeapRegion* form a heap.

Heapsort -- Trace

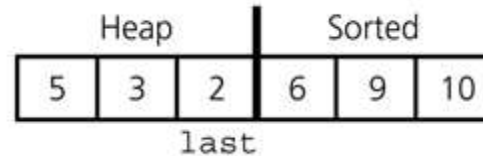


Heapsort -- Trace

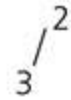
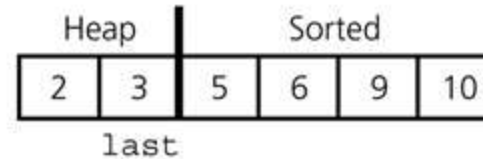
After swapping `anArray[0]` with `anArray[Last]` and decrementing `last`



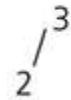
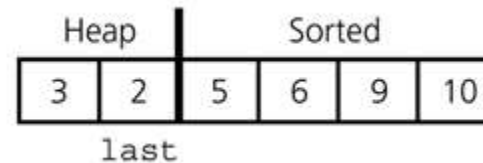
After `rebuildHeap(anArray, 0, 2)`



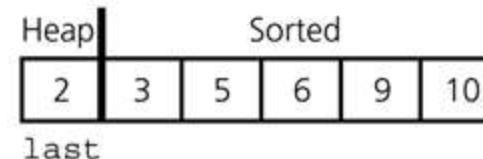
After swapping `anArray[0]` with `anArray[last]` and decrementing `last`



After `heapRebuild(anArray, 0, 1)`



After swapping `anArray[0]` with `anArray[last]` and decrementing `last`



Array is sorted

Heapsort -- Analysis

- Heapsort is
 - $O(n \log n)$ at the average case
 - $O(n \log n)$ at the worst case
- Compared against *quicksort*,
 - *Heapsort* usually takes more time at the average case
 - But its worst case is also $O(n \log n)$.