

2-3 Trees & Red-Black Trees


Initially prepared by Dr. İlyas Çiçekli; improved by various Bilkent CS202 instructors.

2-3 Trees

Definition:

A 2-3 tree is a tree in which each internal node has either two or three children, and all leaves are at the same level.

- **2-node:** a node with two children
- **3-node:** a node with three children



An example of a 2-3 tree

no 1 or 0 child allowed

grows/shrink from the root side

→ A 2-3 tree is not a binary tree

→ A 2-3 tree is never taller than a minimum-height binary tree


→ A 2-3 tree with N nodes never has height greater than $\lceil \log_2(N+1) \rceil$

→ A 2-3 tree of height h always has at least $2^h - 1$ nodes.

2-3 Trees


T is a 2-3 tree of height h if

1. T is empty (a 2-3 tree of height 0), or
2. T is of the form




where r is a node that contains one data item and T_L and T_R are both 2-3 trees, each of height h-1, or

3. T is of the form




where r is a node that contains two data items and T_L , T_M and T_R are 2-3 trees, each of height h-1.

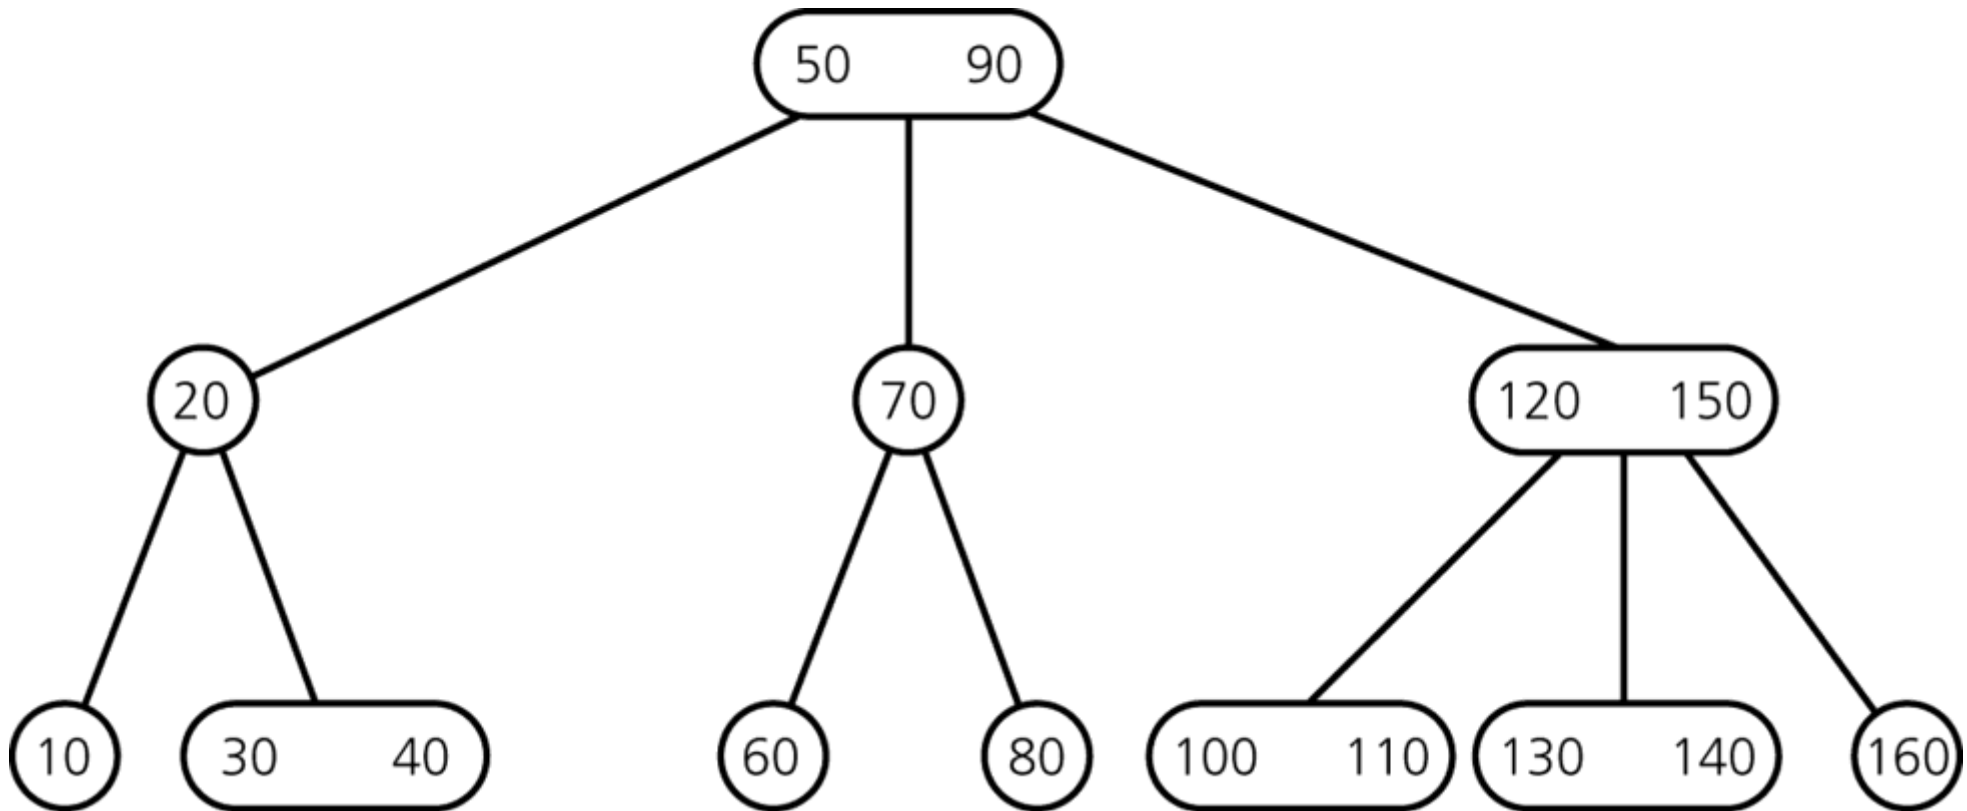
2-node



3-node



2-3 Trees -- Example



C++ Class for a 2-3 Tree Node


```
classTreeNode {  
private:  
    TreeItemTypesmallItem, largeItem;  
    TreeNode *leftChildPtr, *midChildPtr, *rightChildPtr;  
  
    // friend class-can access private class members  
    friend classTwoThreeTree;  
};
```

extra storage

- When a node is a **2-node** (contains only one item)
 - Place it in **smallItem**
 - Use **leftChildPtr** and **midChildPtr** to point to the node's children
 - Place **NULL** in **rightChildPtr**


Traversing a 2-3 Tree

- Inorder traversal visits the nodes in a sorted search-key order
 - Leaf node:
 - Visit the data item(s)
 - 2-node:
 - Visit its left subtree
 - Visit the data item
 - Visit its right subtree
 - 3-node:
 - Visit its left subtree
 - Visit the **smaller data item**
 - Visit its middle subtree
 - Visit the **larger data item**
 - Visit its right subtree




Searching a 2-3 Tree

- Searching a 2-3 tree is similar to searching a binary search tree
 - For a 3-node, compare the searched key with the two values of the 3-node and select one of its three subtrees according to these comparisons
- Searching a 2-3 tree is $O(\log N)$
 - Searching a 2-3 tree and the shortest BST has approximately the same efficiency.
 - A binary search tree with N nodes cannot be shorter than $\lceil \log_2(N+1) \rceil$
 - A 2-3 tree with N nodes cannot be taller than $\lceil \log_2(N+1) \rceil$





A 2-3 tree with the same elements



A balanced binary search tree

Inserting into a 2-3 Tree






Insert [39 38 37 36 35 34 33 32]
into the trees given in the previous slide

- While we insert items into a 2-3 tree, its shape is maintained

Inserting into a 2-3 Tree -- Example

Starting from the following tree, insert [39 38 37 36 35 34 33 32]




Insert 39

- Find the node into which you can put 39


**Insertion
into a 2-node
leaf is simple**

Inserting into a 2-3 Tree -- Example




Insert 38

- Find the node into which you can put 38



Insertion into
a 3-node causes
it to divide

Inserting into a 2-3 Tree -- Example




Insert 37

- Find the node into which you can put 37

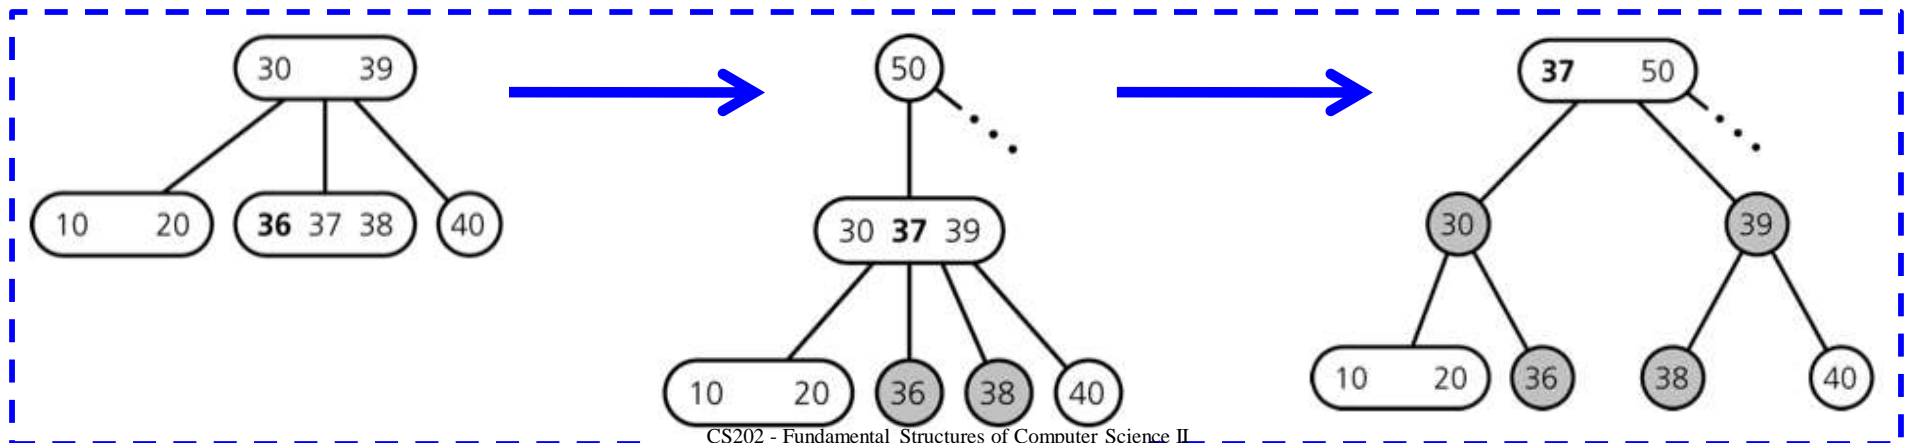
**Insertion
into a 2-node
leaf is simple**

Inserting into a 2-3 Tree -- Example




Insert 36

- Find the node into which you can put 36





Inserting into a 2-3 Tree -- Example



Insert 35, 34, 33



2-3 Trees -- Insertion Algorithm

Splitting a leaf in a 2-3 tree




2-3 Trees -- Insertion Algorithm

Splitting an internal node in a 2-3 tree



2-3 Trees -- Insertion Algorithm

Splitting the root of a 2-3 tree




2-3 tree grows from the root side


Deleting from a 2-3 tree

- Deletion strategy is the inverse of insertion strategy.
- Deletion starts like normal BST deletion (swap with inorder successor)
- Then, we merge the nodes that have become underloaded.

Delete [70 100 80]
from this 2-3 tree




Deleting from a 2-3 Tree -- Example



Delete 70

- Swap with inorder successor
- Delete value from leaf
- Delete the empty leaf
- Shrink the parent (no more mid-pointer)


Deleting from a 2-3 Tree -- Example



Delete 100

- Delete value from leaf
- Distribute the children → Doesn't work
- Redistribute the parent and the children

Deleting from a 2-3 Tree -- Example



Delete 80

- Swap with inorder successor
- Delete value from leaf
- Merge by moving 90 down and removing the empty leaf
- Merge by moving 50 down, adopting empty node's child and removing the empty node
- Remove empty root


all leaves should be on same level

2-3 Trees -- Deletion Algorithm

- To delete an item X from a 2-3 tree:
 - First, we locate the node n containing X.
 - If n is not a leaf, we find X's inorder successor and swap it with X.
 - After the swap, the deletion always begins at the leaf.
 - If the leaf contains another item in addition to X, we simply delete X from that leaf, and we are done.
 - If the leaf contains only X, deleting X would leave the leaf without a data item. In this case, we must perform some additional work to complete the deletion.
- Depending on the empty node and its siblings, we perform certain operations:
 - Delete empty root
 - Merge nodes
 - Redistribute values
- These operations can be repeated all the way upto the root if necessary.

2-3 Trees -- Deletion Operations


Deleting the root




2-3 Trees -- Deletion Operations

Redistributing values (and children)

For a leaf




For an internal node




2-3 Trees -- Deletion Operations

Merging

For a leaf



For an internal node




2-3 Trees -- Analysis

- We can use a 2-3 tree in the implementation of tables.
- A 2-3 tree has the advantage of always being balanced.
- Thus, insertion and deletion operations are $O(\log N)$
- Retrieval based on key is also guaranteed to $O(\log N)$

2-3-4 Trees

- A 2-3-4 tree is like a 2-3 tree, but it allows 4-nodes, which are nodes that have four children and three data items.
- There is a close relation between 2-3-4 trees and red-black trees.
 - We will look at those a bit later
- 2-3-4 trees are also known as 2-4 trees in other books.
 - A specialization of M-way tree ($M=4$)
 - Sometimes also called 4th order B-trees
 - Variants of B-trees are very useful in databases and file systems
 - MySQL, Oracle, MS SQL all use B+ trees for indexing
 - Many file systems (NTFS, Ext2FS etc.) use B+ trees for indexing metadata (file size, date etc.)
- Although a 2-3-4 tree has more efficient insertion and deletion operations than a 2-3 tree, a 2-3-4 tree has greater storage requirements.

2-3-4 Trees -- Example




2-3-4 Trees

T is a 2-3-4 tree of height h if


1. T is empty (a 2-3-4 tree of height 0), or

1. T is of the form




where r is a node containing one data item and T_L and T_R are both 2-3-4 trees, each of height h-1, or

3. T is of the form




where r is a node containing two data items and T_L , T_M and T_R are 2-3-4 trees, each of height h-1, or

4. T is of the form




where r is a node containing three data items and T_L , T_{ML} , T_{MR} , and T_R are 2-3-4 trees, each of height h-1.


2-node



3-node



4-node



C++ Class for a 2-3-4 Tree Node

```
class TreeNode {  
private:  
  
    TreeItemTypes smallItem, middleItem, largeItem;  
  
    TreeNode *leftChildPtr, *lMidChildPtr;  
    TreeNode *rMidChildPtr, *rightChildPtr;  
  
friend class TwoThreeFourTree;  
};
```

- When a node is a 3-node (contains only two items)
 - Place the items in `smallItem` and `middleItem`
 - Use `leftChildPtr`, `lMidChildPtr`, `rMidChildPtr` to point to the node's children
 - Place `NULL` in `rightChildPtr`
- When a node is a 2-node (contains only one item)
 - Place the item in `smallItem`
 - Use `leftChildPtr`, `lMidChildPtr` to point to the node's children
 - Place `NULL` in `rMidChildPtr` and `rightChildPtr`


2-3-4 Trees -- Operations

- Searching and traversal algorithms for a 2-3-4 tree are similar to the 2-3 algorithms.
- For a 2-3-4 tree, insertion and deletion algorithms that are used for 2-3 trees, can similarly be used.
- But, we can also use a slightly different insertion and deletion algorithms for 2-3-4 trees to gain some efficiency.


Inserting into a 2-3-4 Tree

- Splits 4-nodes by moving one of its items up to its parent node.
- For a 2-3 tree, the insertion algorithm traces a path from the root to a leaf and then backs up from the leaf as it splits nodes.
- *To avoid this return path after reaching a leaf*, the insertion algorithm for a 2-3-4 tree splits 4-nodes as soon as it encounters them on the way down the tree from the root to a leaf.
 - As a result, when a 4-node is split and an item is moved up to node's parent, the parent cannot possibly be a 4-node and so can accommodate another item.

Insert[20 50 40 70 80 15 90 100] to
this 2-3-4 tree




Inserting into a 2-3-4 Tree -- Example



Insert 20

- Root is a 4-node → **Split 4-nodes** as they are encountered
- So, we split it before insertion
- And, then add 20


Inserting into a 2-3-4 Tree -- Example



Insert 50 and 40

- No 4-nodes have been encountered → **No split operation** during their insertion


Inserting into a 2-3-4 Tree -- Example



Insert 70

- A 4-node is encountered
- So, we split it before insertion
- And, then add 70


Inserting into a 2-3-4 Tree -- Example



Insert 80 and 15

- No 4-nodes have been encountered → **No split operation** during their insertion


Inserting into a 2-3-4 Tree -- Example



Insert 90

- A 4-node is encountered
- So, we split it before insertion
- And, then add 90

Inserting into a 2-3-4 Tree -- Example



Insert 100


- A 4-node is encountered root
- So, we split it before insertion
- And, then add 100

Splitting 4-nodes during insertion

- We split each 4-node as soon as we encounter it during our search from the root to a leaf that will accommodate the new item to be inserted.
- The 4-node which will be split can:
 - be the root, or
 - have a 2-node parent, or
 - have a 3-node parent.



Splitting 4-nodes during insertion

Splitting a 4-node root






Splitting 4-nodes during insertion

Splitting a 4-node whose parent is a 2-node



Splitting 4-nodes during insertion

Splitting a 4-node whose parent is a 3-node



Deleting from a 2-3-4 tree

- For a 2-3 tree, the deletion algorithm traces a path from the root to a leaf and then backs up from the leaf, fixing empty nodes on the path back up to root.
- *To avoid this return path after reaching a leaf*, the deletion algorithm for a 2-3-4 tree transforms each 2-node into either 3-node or 4-node as soon as it encounters them on the way down the tree from the root to a leaf.
 - If an adjacent sibling is a 3-node or 4-node, transfer an item from that sibling to our 2-node.
 - If adjacent sibling is a 2-node, merge them.



Red-Black Trees

- In general, a 2-3-4 tree requires more storage than a binary search tree.
- A special binary search tree, the **red-black-tree**, can be used to represent a 2-3-4 tree, so that we can retain advantages of a 2-3-4 tree without a storage overhead.
 - 3-node and 4-nodes in a 2-3-4 tree are represented by a binary tree.
 - To distinguish the original 2-nodes from 2-nodes that are generated from 3-nodes and 4-nodes, we use red and black pointers.
 - All original pointers in a 2-3-4 tree are black pointers, red pointers are used for child pointers to link 2-nodes that result from the split of 3-nodes and 4-nodes.



Red-Black Trees

Red-black tree representation


For a 4-node



For a 3-node



or





-- Red pointer
— Black pointer

Red-Black Trees -- Properties

- Root is always a **black node**.
- The children of a **red node** (*pointed by a red pointer*) are always **black nodes** (*pointed by a black pointer*)
- All external nodes (leaves and nodes with a single child) should have **the same number of black pointers** on the path from the root to that external node.

perfect balance

A 2-3-4 Tree and Its Corresponding Red-Black Tree



C++ Class for a Red-Black Tree Node

```
enum Color {RED, BLACK};
```

```
classTreeNode {
```

```
private:
```

```
    TreeItemType Item;
```

```
    TreeNode *leftChildPtr, *rightChildPtr;
```

```
    Color      leftColor, rightColor;
```



only one color variable for node, no red or black pointers is ok too

```
friendclassRedBlackTree;
```

```
};
```





Splitting in a Red-Black Tree Representation

For a 4-node that is the root




Splitting in a Red-Black Tree Representation

For a 4-node whose parent is a 2-node




Splitting in a Red-Black Tree Representation

For a 4-node whose parent is a 3-node



For a 4-node whose parent is a 3-node



Splitting in a Red-Black Tree Representation

For a 4-node whose parent is a 3-node

