

Graphs

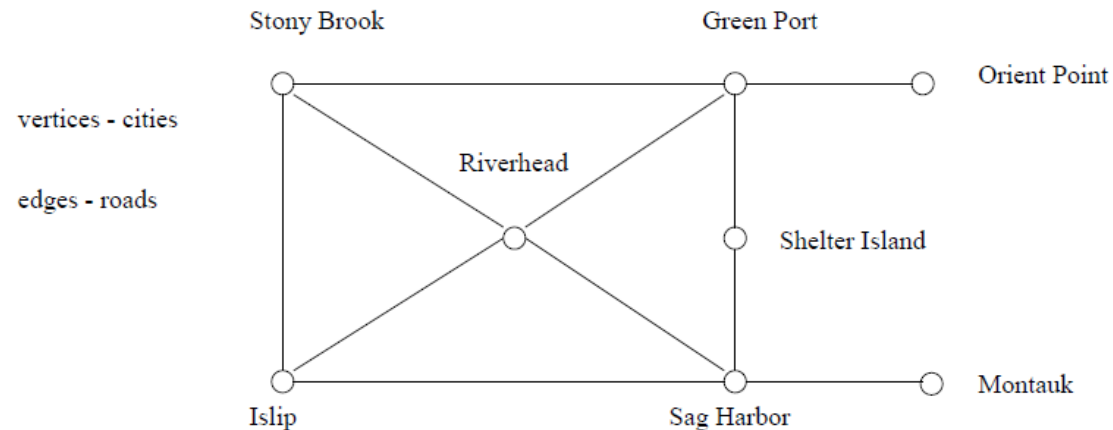
Initially prepared by Dr. İlyas Çiçekli; improved by various Bilkent CS202 instructors.

Graphs

- Graphs are one of the unifying themes of computer science.
- A graph $G = (V, E)$ is defined by a set of *vertices* V , and a set of *edges* E consisting of ordered or unordered pairs of vertices from V .


Road Networks

- In modeling a road network, the vertices may represent the cities or junctions, certain pairs of which are connected by roads/edges.



Electronic Circuits

- In an electronic circuit, with junctions as vertices and components as edges.



Applications

- Social networks (facebook ...)
- Courses with prerequisites
- Computer networks
- Google maps
- Airline flight schedules
- Computer games
- WWW documents
- ... (so many to list!)


Graphs – Definitions

- A **graph** $G = (V, E)$ consists of
 - a set of **vertices**, V , and
 - a set of **edges**, E , where each edge is a pair (v, w) s.t. $v, w \in V$
- Vertices are sometimes called **nodes**, edges are sometimes called **arcs**.
- If the edge pair is ordered then the graph is called a **directed graph** (also called *digraphs*).
- We also call a normal graph (which is not a directed graph) an **undirected graph**.
 - When we say graph we mean that it is an undirected graph.

Graphs – Definitions

- Two vertices of a graph are **adjacent** if they are joined by an edge.
- Vertex w is **adjacent to** v iff $(v,w) \in E$.
 - In an undirected graph with edge (v, w) and hence (w,v) w is adjacent to v and v is adjacent to w .
- A **path** between two vertices is a sequence of edges that begins at one vertex and ends at another vertex.
 - i.e. w_1, w_2, \dots, w_N is a path if $(w_i, w_{i+1}) \in E$ for $1 \leq i \leq N-1$
 - A **simple path** passes through a vertex **only once**.
- A **cycle** is a path that begins and ends at the **same vertex**.
 - A **simple cycle** is a cycle that does not pass through other vertices more than once.

Graphs – An Example



The graph $G = (V, E)$ has 5 vertices and 6 edges:


$$V = \{1, 2, 3, 4, 5\}$$

$$E = \{ (1,2), (1,3), (1,4), (2,5), (3,4), (4,5), (2,1), (3,1), (4,1), (5,2), (4,3), (5,4) \}$$


- *Adjacent:*
1 and 2 are adjacent -- 1 is adjacent to 2 and 2 is adjacent to 1
- *Path:*
1,2,5 (a simple path), 1,3,4,1,2,5 (a path but not a simple path)
- *Cycle:*
1,3,4,1 (a simple cycle), 1,3,4,1,4,1 (cycle, but not simple cycle)

Graphs – Definitions


- A **connected graph** has a path between each pair of distinct vertices.
- A **complete graph** has an edge between each pair of distinct vertices.
 - A complete graph is also a connected graph. But a connected graph may not be a complete graph.



(a) **connected**



(b) **disconnected**



(c) **complete**

Directed Graphs


- If edge pair is ordered, then graph is called a **directed graph** (also called *digraphs*) .
 - Each edge in a directed graph has a direction. Called a **directed edge**.
- Definitions given for undirected graphs apply also to directed graphs, with changes that account for direction.
- Vertex w is **adjacent to** v iff $(v,w) \in E$.
 - i.e. There is a direct edge from v to w
 - w is **successor** of v
 - v is **predecessor** of w
- A **directed path** between two vertices is a **sequence of directed edges** that begins at one vertex and ends at another vertex.
 - i.e. w_1, w_2, \dots, w_N is a path if $(w_i, w_{i+1}) \in E$ for $1 \leq i \leq N-1$

Directed Graphs

- A **cycle** in a directed graph is a path of length ≥ 1 such that $w_1 = w_N$.
 - This cycle is simple if the path is simple.
 - For undirected graphs, the edges must be distinct
- A **directed acyclic graph** (*DAG*) is a directed graph with **no cycles**.
- An undirected graph is **connected** if **there is a path from every vertex to every other vertex**.
 - A directed graph with this property is called **strongly connected**.
 - If a **directed graph** is not strongly connected, but the underlying graph (**without direction to arcs**) is connected then the graph is **weakly connected**.

total cycles --> traverse the graph with a node, if visit the same node, cycle + 1
continue with other nodes

Directed Graph – An Example



The graph $G = (V, E)$ has 5 vertices and 6 edges:



$$V = \{1, 2, 3, 4, 5\}$$

$$E = \{ (1, 2), (1, 4), (2, 5), (4, 5), (3, 1), (4, 3) \}$$

- *Adjacent:*
2 is adjacent to 1, but 1 is NOT adjacent to 2
- *Path:*
1, 2, 5 (a directed path),
- *Cycle:*
1, 4, 3, 1 (a directed cycle),

Weighted Graph

- We can label the edges of a graph with numeric values, the graph is called a **weighted graph**.



Graph Implementations


- The two most common implementations of a graph are:
 - *Adjacency Matrix*
 - A two dimensional array
 - *Adjacency List*
 - For each vertex we keep a list of adjacent vertices

Adjacency Matrix

- An **adjacency matrix** for a graph with n vertices is an n by n array *matrix* such that
 - $matrix[i][j]$ is 1 (true) if there is an edge from vertex i to vertex j
 - 0 (false) otherwise.
- When the graph is **weighted**:
 - $matrix[i][j]$ is weight that labels edge from vertex i to vertex j instead of simply 1,
 - $matrix[i][j]$ equals to ∞ instead of 0 when there is no edge from vertex i to j
- Adjacency matrix for an undirected graph is symmetrical.
 - i.e. $matrix[i][j]$ is equal to $matrix[j][i]$ shortest path etc. can be confusing if 0 is used
- Space requirement $O(|V|^2)$
 - Acceptable if the graph is dense.

Adjacency Matrix – Example1

A directed graph




Its adjacency matrix

		0	1	2	3	4	5	6	7	8
		P	Q	R	S	T	W	X	Y	Z
0	P	0	0	1	0	0	1	0	0	0
1	Q	0	0	0	0	0	0	1	0	0
2	R	0	0	0	0	0	0	1	0	0
3	S	0	0	0	0	1	0	0	0	0
4	T	0	0	0	0	0	1	0	0	0
5	W	0	0	0	1	0	0	0	1	0
6	X	0	0	0	0	0	0	0	0	0
7	Y	0	0	1	0	0	0	0	0	1
8	Z	0	0	0	0	0	0	0	0	0

Adjacency Matrix – Example2

An Undirected Weighted Graph



Its Adjacency Matrix


		0	1	2	3
		A	B	C	D
0	A	∞	8	∞	6
1	B	8	∞	9	∞
2	C	∞	9	∞	∞
3	D	6	∞	∞	∞

Adjacency List

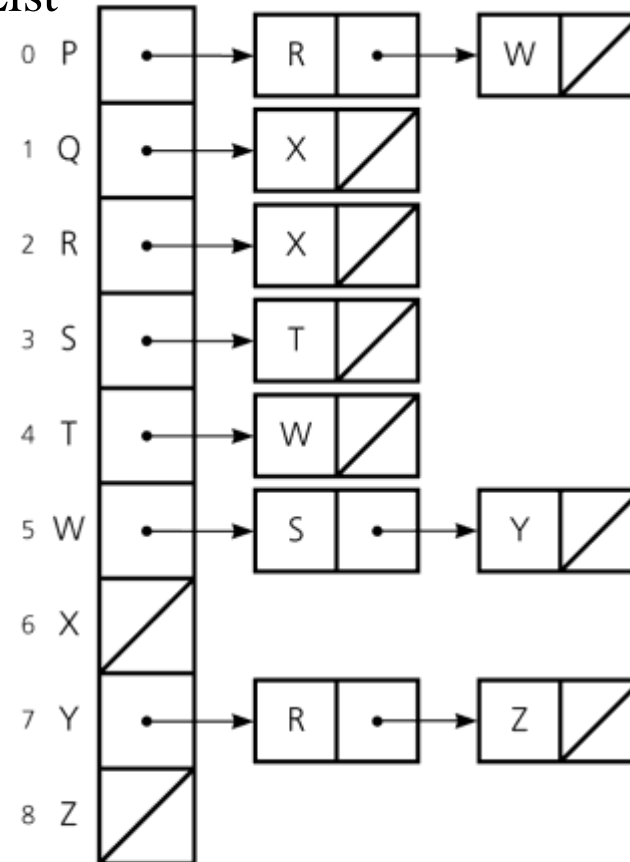
- An **adjacency list** for a graph with n vertices numbered $0, 1, \dots, n-1$ consists of n linked lists. The i^{th} linked list has a node for vertex j if and only if the graph contains an edge from vertex i to vertex j .
- Adjacency list is a better solution if the graph is sparse.
- Space requirement is $O(|E| + |V|)$, which is linear in the size of the graph.
- In an undirected graph each edge (v, w) appears in two lists.
 - Space requirement is doubled.

Adjacency List – Example1

A directed graph




Its Adjacency List




Adjacency List – Example2

An Undirected Weighted Graph



Its Adjacency List



Adjacency Matrix vs Adjacency List

- Two common graph operations:
 1. **Determine whether there is an edge** from vertex i to vertex j .
 2. **Find all vertices adjacent** to a given vertex i .
- An adjacency matrix supports operation 1 more efficiently.
- An adjacency list supports operation 2 more efficiently.
- An adjacency list often requires less space than an adjacency matrix.
 - Adjacency Matrix: Space requirement is $O(|V|^2)$
 - Adjacency List : Space requirement is $O(|E| + |V|)$, which is linear in the size of the graph.
 - Adjacency matrix is better if the graph is dense (too many edges)
 - Adjacency list is better if the graph is sparse (few edges)

Tradeoffs Between Adjacency Lists and Adjacency Matrices

- Faster to test if $(x; y)$ exists?
- Faster to find vertex degree?
- Less memory on sparse graphs?
- Less memory on dense graphs?
- Edge insertion or deletion?
- Faster to traverse the graph?
- Better for most problems?

Tradeoffs Between Adjacency Lists and Adjacency Matrices

- Faster to test if $(x; y)$ exists? **matrices**
- Faster to find vertex degree? **lists**
- Less memory on sparse graphs? **lists – $(m + n)$ vs. (n^2)**
- Less memory on dense graphs? **matrices (small win)**
- Edge insertion or deletion? **matrices $O(1)$**
- Faster to traverse the graph? **lists – $(m + n)$ vs. (n^2)**
- Better for most problems? **lists**

list $\rightarrow O(V+E)$
matrix $\rightarrow O(V^2)$


Graph Traversals

- A **graph-traversal** algorithm starts from a vertex v , **visits all of the vertices** that can be reachable from the vertex v .
 - A graph-traversal algorithm visits all vertices if and only if the graph is **connected**.
- A **connected component** is the subset of vertices visited during a traversal algorithm that begins at a given vertex.
- A graph-traversal algorithm must **mark each vertex** during a visit and **must never visit a vertex more than once**.
 - Thus, if a graph contains a cycle, the graph-traversal algorithm can **avoid infinite loop**.
- We look at two graph-traversal algorithms:
 - **Depth-First Traversal**
 - **Breadth-First Traversal**

Depth-First Traversal

- For a given vertex v , **depth-first traversal** algorithm proceeds along a path from v as deeply into the graph as possible before backing up.
- That is, after visiting a vertex v , the algorithm visits (if possible) an unvisited adjacent vertex to vertex v .
- The depth-first traversal algorithm does not completely specify the order in which it should visit the vertices adjacent to v .
 - We may visit the vertices adjacent to v in sorted order.

Depth-First Traversal – Example



- A depth-first traversal of the graph starting from vertex v.
- Visit a vertex, then visit a vertex adjacent to that vertex.
- If there is no unvisited vertex adjacent to visited vertex, back up to the previous step.


Recursive Depth-First Traversal Algorithm

```
dft(in v:Vertex) {  
    // Traverses a graph beginning at vertex v  
    // by using depth-first strategy  
    // Recursive Version  
    Mark v as visited;  
    for (each unvisited vertex u adjacent to v)  
        dft(u)  
}
```


Iterative Depth-First Traversal Algorithm

```
dft(in v:Vertex) {  
    // Traverses a graph beginning at vertex v  
    // by using depth-first strategy: Iterative Version  
    s.createStack();  
    // push v into the stack and mark it  
    s.push(v);  
    Mark v as visited;  
    while (!s.isEmpty()) {  
        if (no unvisited vertices are adjacent to the vertex on  
            the top of stack)  
            s.pop(); // backtrack  
        else {  
            Select an unvisited vertex u adjacent to the vertex  
                on the top of the stack;  
            s.push(u);  
            Mark u as visited;  
        }  
    }  
}
```

Trace of Iterative DFT – starting from vertex a



Trace of Iterative DFT – starting from vertex a



Node visited


Stack (bottom to top)

a	a
b	a b
c	a b c
d	a b c d
g	a b c d g
e	a b c d g e
(backtrack)	a b c d g
f	a b c d g f
(backtrack)	a b c d g
(backtrack)	a b c d
h	a b c d h
(backtrack)	a b c d
(backtrack)	a b c
(backtrack)	a b
(backtrack)	a
i	a i
(backtrack)	a
(backtrack)	(empty)

Breadth-First Traversal

- After visiting a given vertex v , the **breadth-first traversal** algorithm visits every vertex adjacent to v that it can before visiting any other vertex.
- The breadth-first traversal algorithm does not completely specify the order in which it should visit the vertices adjacent to v .
 - We may visit the vertices adjacent to v in sorted order.

Breadth-First Traversal – Example




- A breadth-first traversal of the graph starting from vertex v.
- Visit a vertex, then visit all vertices adjacent to that vertex.


Iterative Breadth-First Traversal Algorithm

```
bft(in v:Vertex) {  
    // Traverses a graph beginning at vertex v  
    // by using breadth-first strategy: Iterative Version  
    q.createQueue();  
    // add v to the queue and mark it  
    Mark v as visited;  
    q.enqueue(v);  
    while (!q.isEmpty()) {  
        q.dequeue(w);  
        for (each unvisited vertex u adjacent to w) {  
            Mark u as visited;  
            q.enqueue(u);  
        }  
    }  
}
```

Trace of Iterative BFT – starting from vertex a



Trace of Iterative BFT – starting from vertex a

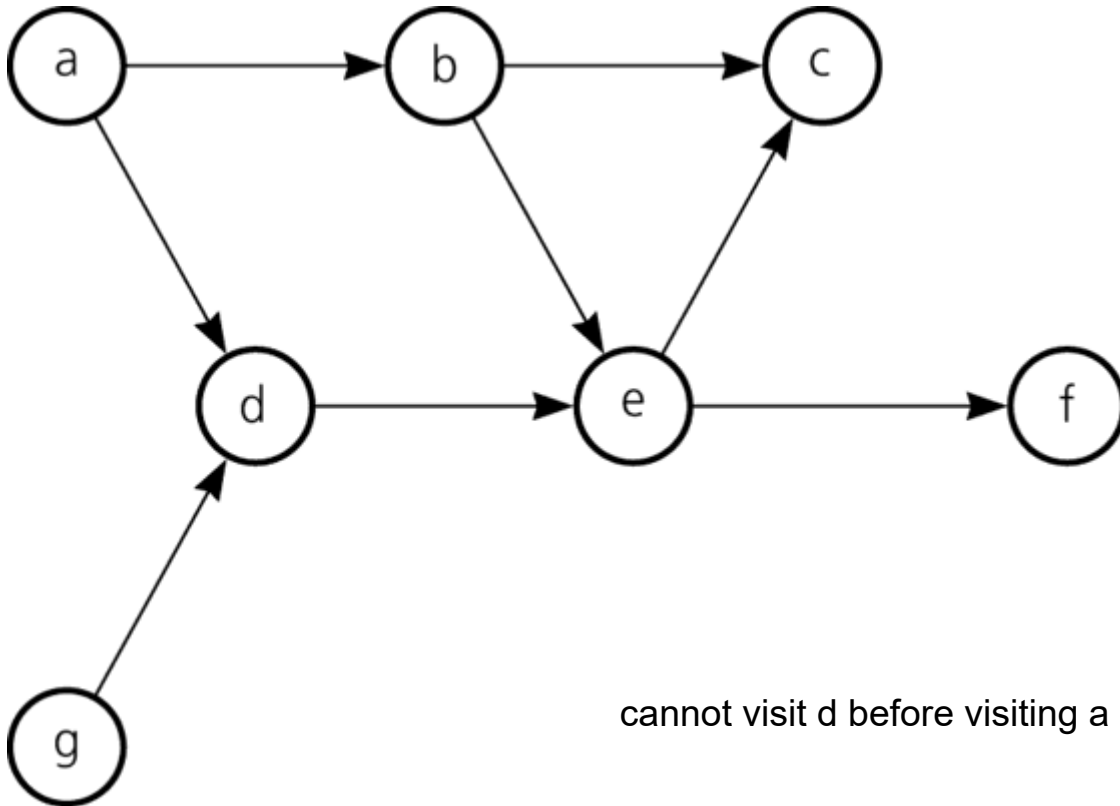


<u>Node visited</u>	<u>Queue (front to back)</u>
a	a
	(empty)
b	b
f	b f
i	b f i
	f i
c	f i c
e	f i c e
	i c e
g	i c e g
	c e g
	e g
d	e g d
	g d
	d
	(empty)
h	h
	(empty)

Topological Sorting

- A directed graph without cycles has a natural order.
 - That is, vertex a precedes vertex b , which precedes c
 - For example, the prerequisite structure for the courses.
- In which order we should visit the vertices of a directed graph without cycles so that we can visit vertex v after we visit its predecessors.
 - This is a linear order, and it is known as **topological order**.
- For a given graph, there may be more than one topological order.
- Arranging the vertices into a topological order is called **topological sorting**.

Topological Order – Example




Some Topological Orders
for this graph:

a, g, d, b, e, c, f

a, b, g, d, e, f, c

cannot visit d before visiting a and g etc.

Topological Order – Example (cont.)




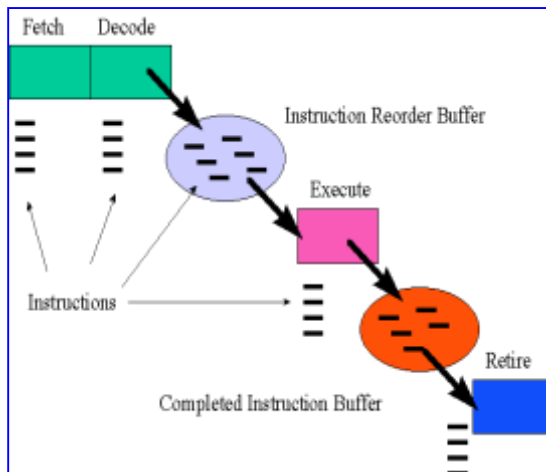


The graph arranged according to the topological orders

(a) *a, g, d, b, e, c, f* and

(b) *a, b, g, d, e, f, c*

Applications of Topological Sorting

- Course prerequisites – which courses should you take next semester?
- Project scheduling (PERT)
- Processors: Instruction scheduling
- Spreadsheets
- Makefiles



Column name Lead value

	y	x	1
1	x	y	
2	1	1	
3	2	4	
4	3	9	
5	4	16	
6	5	25	

Written column

VBA Code

```


Option Explicit
Sub sq()
  Dim n As Integer
  Dim j As Integer
  Dim z As Variant
  'Count size of array
  n = Range("x").Count
  'Do loop
  For j = 1 To n
    z = Range("x").Cells(j).Value
    Range("y").Cells(j).Value = z * z
  Next j
End Sub
  
```

Simple Topological Sorting Algorithm1 – topSort1


```
topSort1(in theGraph:Graph, out aList:List) {  
  // Arranges the vertices in graph theGraph into a  
  // topological order and places them in list aList  
  n = number of vertices in theGraph;  
  for (step=1 through n) {  
    select a vertex v that has no successors;  
    aList.insert(1,v);  
    Delete from theGraph vertex v and its edges;  
  }  
}
```

starting from last vertex f


Graph theGraph




Remove f from theGraph;
add it to aList



Remove c from theGraph;
add it to aList



Remove e from theGraph;
add it to aList



List aList


f

c f

e c f

Graph theGraph

Remove b from theGraph;
add it to aList



Remove d from theGraph;
add it to aList



Remove g from theGraph;
add it to aList



Remove a from theGraph;
add it to aList

List aList

b e c f

d b e c f

g d b e c f

a g d b e c f


DFS Topological Sorting Algorithm – topSort2

```
topSort2(in theGraph:Graph, out aList:List) {  
    // Arranges the vertices in graph theGraph into a topological order and  
    // places them in list aList  
    s.createStack();  
    for (all vertices v in the graph)  
        if (v has no predecessors) {  
            s.push(v);           push a and g (initial)  
                                for starting vertices  
            Mark v as visited;  
        }  
    while (!s.isEmpty()) {  
        if (all vertices adjacent to the vertex on the top of stack  
            have been visited) {  
            s.pop(v);  
            aList.insert(1,v);  
        }  
        else {  
            Select an unvisited vertex u adjacent to the vertex on  
                the top of the stack;  
            s.push(u);  
            Mark u as visited;  
        }  
    }  
}
```

Remember: Iterative Depth-First Traversal Algorithm

```
dft(in v:Vertex) {  
    // Traverses a graph beginning at vertex v  
    // by using depth-first strategy: Iterative Version  
    s.createStack();  
    // push v into the stack and mark it  
    s.push(v);  
    Mark v as visited;  
    while (!s.isEmpty()) {  
        if (no unvisited vertices are adjacent to the vertex on  
            the top of stack)  
            s.pop(); // backtrack  
        else {  
            Select an unvisited vertex u adjacent to the vertex  
                on the top of the stack;  
            s.push(u);  
            Mark u as visited;  
        }  
    }  
}
```


Trace of topSort2



Trace of topSort2

<u>Action</u>	<u>Stack s (bottom to top)</u>	<u>List aList (beginning to end)</u>
Push a	a	
Push g	a g	
Push d	a g d	
Push e	a g d e	
Push c	a g d e c	
Pop c, add c to aList	a g d e	c
Push f	a g d e f	c
Pop f, add f to aList	a g d e	f c
Pop e, add e to aList	a g d	e f c
Pop d, add d to aList	a g	d e f c
Pop g, add g to aList	a	g d e f c
Push b	a b	g d e f c
Pop b, add b to aList	a	b g d e f c
Pop a, add a to aList	(empty)	a b g d e f c

top = g
continue with unvisited vertices
adjacent to g, then d, then e etc.
c and f both adjacent to e, use sorted order




top = a, only adjacent unvisited is b, push b

Spanning Trees

- A tree is a special kind of undirected graph.
- That is, a **tree** is a connected undirected graph without cycles.
- All trees are graphs, not all graphs are trees. **Why?**
- A **spanning tree** of a connected undirected graph G is a sub-graph of G that contains all of G 's vertices and enough of its edges to form a tree.
- There may be several spanning trees for a given graph.
- If we have a connected undirected graph with cycles, and we remove edges until there are no cycles to obtain a spanning tree.

A Spanning Tree



Remove dashed lines to obtain a spanning tree




Cycles?

Observations about graphs:

reverse is not true

1. A connected undirected graph that has n vertices must have at least $n-1$ edges.
2. A connected undirected graph that has n vertices and exactly $n-1$ edges cannot contain a cycle.
(tree)
3. A connected undirected graph that has n vertices and more than $n-1$ edges must contain a cycle.

cycle = minimum 2 edges



Connected graphs that each have four vertices and three edges


DFS Spanning Tree

```
dfsTree (in v:vertex) {  
    // Forms a spanning tree for a connected undirected graph  
    // beginning at vertex v by using depth-first search;  
    // Recursive Version  
    Mark v as visited;  
    for (each unvisited vertex u adjacent to v) {  
        Mark the edge from u to v;  
        dfsTree (u);  
    }  
}
```

Remember: Recursive Depth-First Traversal Algorithm

```
dft(in v:Vertex) {  
    // Traverses a graph beginning at vertex v  
    // by using depth-first strategy  
    // Recursive Version  
    Mark v as visited;  
    for (each unvisited vertex u adjacent to v)  
        dft(u)  
}
```

DFS Spanning Tree – Example



The DFS spanning tree algorithm visits vertices in this order: a, b, c, d, g, e, f, h, i. Numbers indicate the order in which the algorithm marks edges.

The DFS spanning tree rooted at vertex *a*


BFS Spanning tree

```
bfsTree(in v:vertex) {  
    // Forms a spanning tree for a connected undirected graph  
    // beginning at vertex v by using breadth-first search;  
    // Iterative Version  
    q.createQueue();  
    q.enqueue(v);  
    Mark v as visited;  
    while (!q.isEmpty()) {  
        q.dequeue(w);  
        for (each unvisited vertex u adjacent to w) {  
            Mark u as visited;  
            Mark edge between w and u;  
            q.enqueue(u);  
        }  
    }  
}
```

Remember: Iterative Breadth-First Traversal Algorithm

```
bft(in v:Vertex) {  
    // Traverses a graph beginning at vertex v  
    // by using breadth-first strategy: Iterative Version  
    q.createQueue();  
    // add v to the queue and mark it  
    q.enqueue(v);  
    Mark v as visited;  
    while (!q.isEmpty()) {  
        q.dequeue(w);  
        for (each unvisited vertex u adjacent to w) {  
            Mark u as visited;  
            q.enqueue(u);  
        }  
    }  
}
```

BFS Spanning tree – Example



The BFS spanning tree algorithm visits vertices in this order: a, b, f, i, c, e, g, d, h. Numbers indicate the order in which the algorithm marks edges.

The BFS spanning tree rooted at vertex *a*

Minimum Spanning Tree

- If we have a weighted connected undirected graph, the edges of each of its spanning tree will also be associated with costs.
- The *cost of a spanning tree* is the sum of the costs of edges in the spanning tree.
- A **minimum spanning tree** of a connected undirected graph has a minimal edge-weight sum.
 - A minimum spanning tree of a connected undirected may not be unique.
 - Although there may be several minimum spanning trees for a particular graph, their costs are equal.


Prim's Algorithm

- **Prim's algorithm** finds a minimum spanning tree that begins any vertex.
- Initially, the tree contains only the starting vertex.
- At each stage, the algorithm selects a **least-cost edge** from among those that begin with a vertex in the tree and end with a vertex not in the tree.
- The selected vertex and least-cost edge are added to the tree.

Prim's Algorithm


```
primsAlgorithm(in v:Vertex) {  
  // Determines a minimum spanning tree for a weighted,  
  // connected, undirected graph whose weights are  
  // nonnegative, beginning with any vertex.  
  Mark vertex v as visited and include it in  
    the minimum spanning tree;  
  while (there are unvisited vertices) {  
    Find the least-cost edge (v,u) from a visited vertex v  
      to some unvisited vertex u;  
    Mark u as visited;  
    Add the vertex u and the edge (v,u) to the minimum  
      spanning tree;  
  }  
}
```

Prim's Algorithm – Trace



A weighted, connected, undirected graph


Prim's Algorithm – Trace (cont.)




(a) Mark a , consider edges from a

beginning at vertex a

Prim's Algorithm – Trace (cont.)




(a) Mark a, consider edges from a





(b) Mark i, include edge (a, i)

beginning at vertex a




Prim's Algorithm – Trace (cont.)







Prim's Algorithm – Trace (cont.)




Prim's Algorithm – Trace (cont.)




Prim's Algorithm – Trace (cont.)




Prim's Algorithm – Trace (cont.)



(g) Mark c, include edge (d, c)



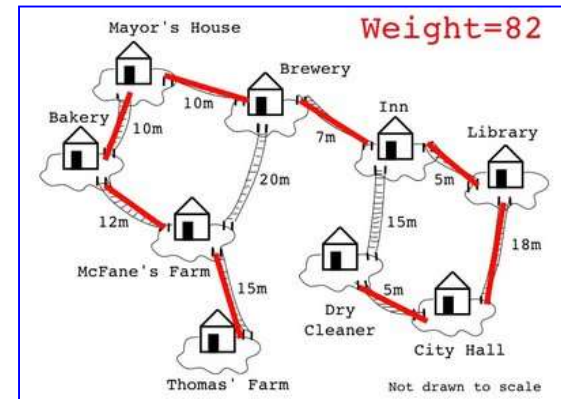


(h) Mark e, include edge (c, e)



(i) Mark b, include edge (a, b)

Applications of Minimum Spanning Trees


- Building low-cost network
- Grouping objects into clusters, where each cluster is a set of similar objects.
- Minimum bottleneck



Shortest Paths

- The **shortest path** between two vertices in a weighted graph has the smallest edge-weight sum.
- **Dijkstra's shortest-path algorithm** finds the shortest paths between vertex 0 (a given vertex) and all other vertices.

Shortest Paths



(b)

	0	1	2	3	4
0	∞	8	∞	9	4
1	∞	∞	1	∞	∞
2	∞	2	∞	3	∞
3	∞	∞	2	∞	7
4	∞	∞	1	∞	∞

A Weighted Directed Graph

Its Adjacency Matrix


Dijkstra's Shortest-Path Algorithm

```
shortestPath(in theGraph, in weight:WeightArray) {  
    // Finds the minimum-cost paths between an origin vertex (vertex 0)  
    // and all other vertices in a weighted directed graph theGraph;  
    // theGraph's weights are nonnegative  
    Create a set vertexSet that contains only vertex 0;  
    n = number of vertices in the Graph;  
    // Step 1  
    for (v=0 through n-1)  
        weight[v] = matrix[0][v];  
    // Steps 2 through n  
    for (v=2 through n) {      //n-1 times v-->i  
        Find the smallest weight[v] such that v is not in vertexSet;  
        Add v to vertexSet;  
        for (all vertices u adjacent to v but not in vertexSet)  
            if (weight[u] > weight[v]+matrix[v][u])  
                weight[u] = weight[v]+matrix[v][u];  
    }  
}
```

Dijkstra's Shortest-Path Algorithm – Trace

Step	<u>v</u>	<u>vertexSet</u>	weight					
			<u>[0]</u>	<u>[1]</u>	<u>[2]</u>	<u>[3]</u>	<u>[4]</u>	
1	–	0	0	8	∞	9	4	min = 4
2	4	0, 4	0	8	5	9	4	
3	2	0, 4, 2	0	7	5	8	4	
4	1	0, 4, 2, 1	0	7	5	8	4	
5	3	0, 4, 2, 1, 3	0	7	5	8	4	

(a) Origin




[2] is adjacent to [4] check
 $\infty > 4 + 1$ update
 [1] and [3] are adjacent to [2] check
 $8 > 5 + 2$ update [1]
 $9 > 5 + 3$ update [3]
 [2] is adjacent to [1]
 check $7 + 1 > 5$ dont update
 etc.

(b)

	0	1	2	3	4
0	∞	8	∞	9	4
1	∞	∞	1	∞	∞
2	∞	2	∞	3	∞
3	∞	∞	2	∞	7
4	∞	∞	1	∞	∞


Dijkstra's Shortest-Path Algorithm – Trace (cont.)

(a)



Step 2. The path 0-4-2 is shorter than 0-2


(b)



Step 3. The path 0-4-2-1 is shorter than 0-1


Dijkstra's Shortest-Path Algorithm – Trace (cont.)

(c)



Step 3 continued. The path 0–4–2–3 is shorter than 0–3

(d)



Step 4. The path 0–4–2–3 is shorter than 0–4–2–1–3